

User Journeys

DCSA Unified Application Access and Integration Prototype (DCSA-UAL)

Field	Value
Product Name	DCSA Unified Application Access and Integration Prototype
Project Acronym	DCSA-UAL
Document Type	User Journeys
Status	Draft v1.0
Date	2026-09-15
Related Personas	project_specs/PERSONAS-DCSA-UAL.md (PER-01 ... PER-04)
Related JTBD	project_specs/JTBD-DCSA-UAL.md (JTBD-01.1 ... JTBD-04.4)
Related PRD	project_specs/PRD-DCSA-UAL.md (§5 Feature Requirements F0-F19; §7 Success Metrics)
Related FRD	project_specs/FRD-DCSA-UAL.md (screen inventory SCR-01 ... SCR-38; seed corpus F17)
Related Charter	.planning/PROJECT.md

DEMO — SYNTHETIC DATA ONLY. Every journey below is written to be **walked live, on synthetic data, in front of an evaluator**. The people are fabricated. The cases, questionnaires, issue items, and identifiers are fabricated and invalid by construction. No real DCSA personnel, applicants, or systems are represented.

How to Read This Document

Each journey is a **to-be** journey: it maps the step the user actually takes in the unified layer, and — in the **Pain Point** column — what that same step costs them **today**, in the fragmented multi-application estate. That column is the product's argument. If a stage has no "today" cost, the stage is not earning its place.

Every journey carries:

- **Entry Trigger** — the real-world event that starts it.
- **Preconditions (Seed Data Required)** — the exact synthetic rows that must exist for the journey to be demonstrable. **This section is normative input to the F17 seeding work.** A journey whose preconditions are not seeded is a journey that cannot be demoed, and seed validation (FR-F17-10) should fail loudly rather than let it be discovered in front of a reviewer.

- **Journey Stages** — step, action, touchpoint (a specific screen from the FRD inventory), the user's thinking, their feeling, today's cost, and the design opportunity.
- **System Response by Stage** — what the product does in reply, with the PRD feature IDs (F0–F19) exercised.
- **Emotion Curve** — the affective shape of the journey on a 1–5 scale (1 = anxious/distrustful, 5 = confident). Journeys that never dip are lies.
- **Key Moments, Success Exit Criteria, Failure and Alternate Paths, Accessibility Notes, Success Outcome, and Feature Touchpoints.**

Accessibility is a completion criterion, not a note. Every journey in this document must be completable **keyboard-only and with a screen reader**, start to finish. Where a journey has a step that is a known accessibility risk — a cross-application traversal, a live connection test, a live-region degraded warning arriving mid-form — the risk is named explicitly in that journey's Accessibility Notes so it can be tested rather than assumed (F14, NFR-01, NFR-02, NFR-16, SM-07, SM-08).

Touchpoint notation. Touchpoints name the FRD screen (e.g. SCR-15 eApp case view) and the component within it, so a journey step maps to something a builder can point at. Screens are defined in FRD-DCSA-UAL.md §Screen Inventory.

Persona ↔ seeded identity binding (action for the seeding work)

The persona names used throughout this document (Marcus Vale, Dana Okonkwo, Renée Ashford, Priya Raghunathan) are the **design personas** from PERSONAS-DCSA-UAL.md. The FRD's seed corpus (FR-F17-02) currently names its synthetic identities differently (Dana R. Whitfield, Marcus T. Okonkwo, Priya L. Raghavan, Corinne A. Vasquez-Hale, Theodore Q. Lansbury, Wilhelmina R. Farraday). **These two name sets must be reconciled before the demo**, in whichever direction the team prefers — but one name per human, everywhere, or the demo script and the persona documents will contradict each other on screen in front of an evaluator. This document's preconditions are written against the **stable seed references** (CASE-A-1042, ISS-2207, SUBJ-00418, SUBJ-00622, CVS), which are unaffected by the naming decision.

Journey Index

ID	Persona	Scenario	Key JTBD	Stages	Demo path
JRN-01.01	PER-01 Marcus Vale	Flagship — clear a blocking PVQ issue on an eApp case in one sitting, one login, dual-system update	JTBD-01.1	11	Primary — Segment 1
JRN-01.02	PER-01 Marcus Vale	Monday-morning triage: one queue answers "what do I work next" across five spokes	JTBD-01.2, JTBD-01.3	7	Supporting — Segment 3b
JRN-01.03	PER-01 Marcus Vale	Keep working while Investigation Management is down — named degradation, disabled actions, automatic recovery	JTBD-01.4	7	Secondary — Segment 2a
JRN-02.01	PER-02 Dana Okonkwo	Assemble the whole cross-system picture on a subject and render a determination	JTBD-02.1, JTBD-02.2	9	Secondary — Segment 4
JRN-02.02	PER-02 Dana Okonkwo	Work the aging queue, return a case for more investigation, and refuse to decide on data that is quietly missing	JTBD-02.3, JTBD-02.4	7	Supporting — Segment 2c
JRN-03.01	PER-03 Renée Ashford	On a phone at lunch: where do I stand, what do I owe, do it now	JTBD-03.1, JTBD-03.2, JTBD-03.3	8	Supporting — Segment 5a

ID	Persona	Scenario	Key JTBD	Stages	Demo path
JRN-03.02	PER-03 Renée Ashford	Prove what I submitted — and see, live, that the boundary around my record holds	JTBD-03.4	6	Secondary — Segment 5b
JRN-04.01	PER-04 Priya Raghunathan	Register the sixth application live, through the UI, and watch it start reporting health	JTBD-04.1	9	Secondary — Segment 3a
JRN-04.02	PER-04 Priya Raghunathan	Triage a reported integration failure: health view → integration-issue log → correlation ID → audit chain → contain	JTBD-04.2, JTBD-04.3, JTBD-04.4	8	Supporting — Segment 6
JRN-04.03	PER-04 Priya Raghunathan	Rehearse degradation on purpose: induce an outage, watch the product degrade visibly, restore it	JTBD-04.2, JTBD-04.4	6	Secondary — Segment 2b

Coverage check. Four personas, ten journeys, every persona ≥ 2 . All sixteen JTBD entries are claimed by at least one journey (see *Journey-to-JTBD Traceability*). Every journey is demonstrable live on the F17 synthetic corpus with no manual setup beyond the documented reset command (FR-F17-11).

The Scripted Demonstration Path

This is the order an evaluator should be walked through, and the order rehearsal should protect. It follows the demo-script assignment already made in JTBD-DCSA-UAL.md (§Downstream Use): **JTBD-01.1 is the primary script; JTBD-01.4, JTBD-02.2, JTBD-03.4, and JTBD-04.1 are the secondary scripts** (PRD F18).

Segment	Journeys	What the evaluator is being shown	Duration
1. The thesis	JRN-01.01	One sign-in, one shell, a case-advancing action that spans eApp and PVQ, and both spokes independently confirming the change. If only one thing is demonstrated, it is this.	~5 min
2. Resilience	JRN-04.03 (drive) → JRN-01.03 (observe) → JRN-02.02 (the adjudicator consequence)	Degradation is named and quantified, not blank; the rest of the queue stays usable; recovery needs no reload and no re-authentication. Two windows side by side: Priya induces, Marcus observes.	~4 min
3. Extensibility	JRN-04.01 (drive) → JRN-01.02 (observe)	A sixth application onboarded as configuration, live, in under five minutes — and its work items appearing in an already-signed-in investigator's queue without him signing out.	~6 min
4. Zero trust, role-level	JRN-02.01	The same PVQ issue work item, opened by the adjudicator and by the investigator, presenting different server-computed action sets .	~3 min
5. Zero trust, resource-level	JRN-03.01 → JRN-03.02	The strictest scope in the product, then the curl denial: another subject's resource ID, denied server-side, non-enumerably, and the denial visible in the audit trail.	~4 min
6. Accountability	JRN-04.02	One correlation ID resolving the entire cross-system chain — including the administrator's own actions, because administrators are not exempt.	~3 min

Rehearsal rule. Segments 2 and 3 mutate platform state (failure injection, registry). The reset command (FR-F17-11) must be run between full rehearsals so Segment 1 is always driven against pristine baseline data — ISS-2207 must be OPEN at the moment the demo starts (SM-22).

PER-01: Marcus Vale

Field/desk background investigator, DCSA Personnel Vetting. 38-45 active cases across IM, eApp, PVQ, and PDT. Measured on cases closed on time and on the defensibility of what he recorded. Interrupted constantly; half his week is on a VPN from somewhere that is not an office.

JRN-01.01: Clear the Blocking Issue — eApp Case → Related PVQ Issue → Dual-System Update

THE FLAGSHIP JOURNEY. This is the single most important journey in this document and the artifact that proves the entire product thesis. It maps PRD **F7** and **JTBD-01.1**. Per the project charter: *if everything else fails, this must work.*

Persona: PER-01 (Marcus Vale) **Scenario:** It is 0815 on a Tuesday. Marcus is at a hoteling desk in the field office with a case that is due in four days and cannot move. eApp case CASE-A-1042 is under review, and PVQ has raised an issue item against one specific answer in its Section 13A employment history — an end date on the subject's first listed employer that does not reconcile with what the employer's HR record says. Marcus has to read the answer, judge it against what he already knows from the employer interview he conducted last week, and dispose of the issue. Today that is two applications, two logins, a scratch document holding three identifiers, and forty minutes. In the unified layer it is one sitting, one authentication, zero re-typed identifiers, and both systems provably updated before he stands up. **Related Jobs:** JTBD-01.1 (*primary*), JTBD-01.2 (*entry*), JTBD-01.3 (*the record he leaves*) **Entry Trigger:** A dashboard alert — ALERT-NEW-PVQ-ISSUE — fires on a case Marcus owns, telling him an issue was raised three days ago against CASE-A-1042. Today, that alert does not exist; he would find out when the case aged or when his supervisor asked. **Demo Path:** □ **PRIMARY SCRIPTED DEMO — Segment 1.** Target: complete manually in **under three minutes** with no instruction beyond the demo script (PRD F7 acceptance signal).

Preconditions (Seed Data Required)

Normative input to F17 seeding. Seed validation (FR-F17-10) must fail loudly if any row below is missing — a broken flagship demo must be diagnosable in seconds, not debugged live.

#	Namespace	Required seed state
P1	hub	An INVESTIGATOR identity bound to PER-01, CAC/PIV-enabled, org=DCSA-FIELD-OPS-EAST , tier=T5 , region=REGION-NE , holding the case assignment for CASE-A-1042 .
P2	eapp	CASE-A-1042 , subject SUBJ-00418 , caseState=UNDER_REVIEW , outstandingIssueCount=1 , outstandingIssueRefs=["ISS-2207"] , due in 4 days, with a populated SECTION_13A containing ≥ 2 employer entries.
P3	pvq	ISS-2207 , status=OPEN , parentSystem=EAPP , parentCaseRef=CASE-A-1042 , subjectRef=SUBJ-00418 , answerLocus=SECTION_13A.employer[0].endDate , answerSectionLabel="Section 13A – Employment history" , answerSnapshot populated, raised 3 days ago.
P4	pdt	A designation referencing CASE-A-1042 , so the related-items panel shows more than one relationship type and the cross-system story is not a single link.
P5	im	An assignment referencing CASE-A-1042 assigned to PER-01, giving the panel a third relationship type.

#	Namespace	Required seed state
P6	hub	PER-01's default queue view surfaces EAPP:CASE-A-1042 on page 1 without filtering ; the ALERT-NEW-PVQ-ISSUE rule fires for ISS-2207 so the dashboard on-ramp is populated.
P7	pvq	A second OPEN issue on a different case owned by PER-01, so the journey can be re-run on alternate data if ISS-2207 has already been consumed mid-session.
P8	hub	~120 pre-seeded audit events so the audit viewer is populated before this journey writes to it, and the new correlated chain is visibly <i>one chain among many</i> rather than the only content.

Journey Stages

Stage	Action	Touchpoint	Thinking	Feeling	Pain Point	Opportunity
1. Sign in	Selects CAC/PIV, picks his synthetic certificate identity	SCR-01 login → SCR-02 certificate picker	"One card, one door. Let's see if it holds."	Neutral, mildly skeptical	Today: three separate logins for one case-advancing action — IM, eApp, PVQ — and the third has always timed out by the time he reaches it	One authentication event for the whole session, asserted in the audit log (SM-02)
2. Orient	Scans the alerts panel and sees "New issue raised on CASE-A-1042 — 3 days ago"	SCR-09 investigator dashboard → alerts panel	"There it is. Nobody had to call me."	Alert, slightly annoyed the issue sat three days	Today: discovery is accidental — he learns a PVQ issue exists when the case is already late or when his supervisor asks	Exceptions are announced and link directly into the item that produced them (F15)
3. Enter the work	Activates the alert link, landing in his queue with the case in view	SCR-13 unified work queue, default view "assigned to me, due date ascending"	"Due in four days. This is the one."	Focused	Today: three lists — IM cases, PVQ issues, eApp questionnaires — none prioritized against the others, so he prioritizes by habit	One queue, five sources, each row badged with its owning system; the next action is ≤ 2 clicks from sign-in
4. Open the case	Opens CASE-A-1042 from the queue row	SCR-15 eApp case view	"Same header, same chrome. I didn't go anywhere."	Settling in	Today: opening the case means a second application and re-establishing in his own head what he was doing	Detail opens inside the shell with queue context preserved for the return trip (F6)
5. Read the answer	Reads Section 13A, employer[0], and the end date under question	SCR-15 → questionnaire section panel, source-badged eApp	"End date says March. The employer told me January."	Engaged, professionally suspicious	Today: the answer is in eApp and the issue is in PVQ, and nothing on his screen connects them	The flagged answer is rendered in the same view that will link to the issue raised against it
6. Discover the relationship	Reads the related-items panel: "Issue raised against Section 13A employment history" plus the PDT designation and IM assignment	SCR-15 → related items panel (live via the PVQ adapter, not a hard-coded link)	"It's already connected. I don't have to go find it."	Relief — this is the moment the product earns him	Today: he copies the issue reference into a scratch document and searches PVQ for it by hand; a transposed digit is a real and recurring error class	Cross-system relationships resolved live and explained in words, with source attribution (F6 → on-ramp to F7)

Stage	Action	Touchpoint	Thinking	Feeling	Pain Point	Opportunity
7. Traverse	Activates the related issue and lands on the PVQ issue, breadcrumb reading <code>Work Queue > eApp Case A-1042 > Related PVQ Issue ISS-2207</code>	SCR-16 PVQ issue detail (in-shell, no interstitial, no new tab)	"No login. No new tab. I'm still in the same place."	Confident	Today: a second authentication, a manual search, and a cognitive reset that costs more than the clicks	Zero re-authentication, zero context re-entry, breadcrumb carrying the case context forward (F1, F3, F7)
8. Resolve	Selects a disposition, types a resolution narrative referencing the employer interview, submits	SCR-16 → resolution form (USWDS, validated, error summary on failure)	"This has to read right to an adjudicator in nine months."	Deliberate, high-stakes concentration	Today: the narrative is written without the questionnaire answer on screen, so he over-explains to compensate for the reader's missing context	Answer and issue in one continuous path; a long narrative is never lost to a validation failure (F14)
9. See both systems answer	Reads the dual-system confirmation naming each spoke and what each now says	SCR-20 dual-system confirmation	"PVQ says resolved. eApp says the case is clear. Both of them said it, not the middle layer."	Trust — the decisive moment	Today: he re-opens each system separately to check the change actually landed, because no confirmation names a system or quantifies anything	State read back from each spoke independently , so the demo proves the dual update rather than asserting it (SM-03)
10. Verify on the case	Returns to the case and sees "No outstanding issues" and the updated case state	SCR-15 re-read from eApp	"Clean. It moves."	Satisfied	Today: verification is a fourth login and a manual comparison of two screens	One click back to a re-read (not cached) case, proving the parent state changed (F7 verification affordance)
11. Leave the record	Opens the item's activity history and sees the whole action as one correlated chain	SCR-15 activity history → SCR-34 audit chain view	"One story, not four rows. That'll hold up."	Reassured, done	Today: correlation is manual — he exports activity from each system and lines timestamps up by hand	One correlation ID spanning both reads, both writes, and the confirmation, viewable as a single narrative (F13, SM-20)

System Response by Stage

Stage	System Response	Features
1. Sign in	Simulated IdP establishes one server-side session; principal (identity, roles, attributes, correlation ID) is now presentable to every adapter; <code>AUTH_SUCCESS</code> audit event written — the only one for the session	F0, F1, F13
2. Orient	Dashboard composes server-side across five spokes; the alert rule computes "new PVQ issue on a case you own" from live spoke data and renders it with a direct deep link	F4, F15, F2
3. Enter the work	Queue fans out to five adapters in one request with per-source partial-failure tolerance; rows normalized and source-badged; result count announced	F5, F8, F14
4. Open the case	Hub authorizes this principal against this case specifically (resource-level, not route-level), fetches detail via the eApp adapter, preserves queue filters/sort/page for return	F2, F6, F8

Stage	System Response	Features
5. Read the answer	Section content rendered from eApp with source attribution present in the accessible name, not just a badge	F6, F14
6. Discover the relationship	Hub resolves outstandingIssueRefs into typed RelatedRef objects by calling the PVQ adapter live; PDT and IM relationships resolved the same way; a reference that will not resolve renders a "couldn't be confirmed" state rather than a broken link	F6, F7, F8, F9
7. Traverse	In-shell route change to the PVQ item carrying from=EAPP:CASE-A-1042 ; no credential prompt, no tab; descriptive page title change; focus placed on the new h1 ; traversal itself audited so the chain shows the path, not just the endpoints	F1, F3, F7, F13, F14
8. Resolve	Server re-authorizes two permissions — ISSUE.RESOLVE on PVQ:ISS-2207 and CASE.UPDATE_ISSUE_STATE on EAPP:CASE-A-1042 — before anything is written; validation errors return an error summary with in-page links and the narrative intact	F2, F6, F7, F14
9. See both systems answer	Hub orchestrates PVQ first (authoritative outcome), then eApp's parent-state update; audit written before the success response returns; if the second write fails, an explicit partial-completion state names exactly which system changed and which did not, with a retry path — never a success message	F7, F13, F10
10. Verify on the case	eApp re-read through its adapter (not from cache) showing outstandingIssueCount=0 ; a per-system view offers PVQ's own record of the resolution	F6, F7, F9
11. Leave the record	Spoke-native history merged with hub audit records into one chronology; chain view returns every record sharing the workflow correlation ID	F6, F13

Emotion curve (1 anxious → 5 confident): Sign in **3** → Orient **3** (an issue sat three days) → Enter **4** → Open case **4** → Read answer **4** → Discover relationship **5** (the "oh — it's already connected" moment) → Traverse **5** → Resolve **3** (deliberate dip: this is a government record about a person, and he is being careful) → Both systems answer **5** (peak trust) → Verify **5** → Record **5**.

Key Moments

- **Decision Point — Stage 8 (Resolve).** Marcus chooses a disposition and writes a narrative that an adjudicator will read months later out of context. This is the product's most consequential write and the one that must never fail silently or lose typed text.
- **Delight Opportunity — Stage 6 (Discover the relationship).** The instant the related-items panel shows the PVQ issue already connected to the answer it was raised against, with the relationship explained in words, the product has replaced his scratch document. This is the moment to slow down in the demo.
- **Peak-trust Moment — Stage 9 (Dual-system confirmation).** Showing each spoke's **own** answer, and being able to prove it with independent API calls outside the hub entirely, is what separates a demonstration from an assertion.
- **Risk of Abandonment — Stage 7 (Traverse).** Any credential prompt, new tab, "you are now leaving this application" interstitial, or visual discontinuity here is the exact pain he came to escape. One occurrence and he stops believing the rest of the product.
- **Trust-destroying Risk — Stage 9.** A success message that turns out to be false in one of the two systems ends this persona's relationship with the product permanently. The partial-completion state exists for this reason and must be demonstrable on demand.

Success Exit Criteria

- The workflow completes end to end in **under three minutes**, manually, following the demo script (PRD F7 acceptance signal, SM-01).
- The audit log contains **exactly one authentication event** for the session (SM-02).

- **Zero** manual re-entry of subject, case, or issue identifiers at any step (SM-04).
- GET {pvq}/issues/ISS-2207 and GET {eapp}/cases/CASE-A-1042 , called **directly against the spokes**, both return the updated state (SM-03).
- The entire action is retrievable as a **single correlated chain** in the audit viewer (SM-20).
- Returning to the queue restores prior filters, sort, and page.
- The whole path is completed **keyboard-only** in a separate verification pass (SM-08).

Failure and Alternate Paths

Condition	What Marcus sees	Recovery
PVQ write succeeds, eApp parent-state write fails	Explicit partial-completion state: "The issue was resolved in PVQ. The case record in eApp was not updated." Named systems, correlation ID, retry action. No success confirmation is shown.	Retry the eApp half; the retry is a distinct audited operation under the same correlation ID
PVQ unavailable at Stage 7	The related issue link renders, but the issue detail shows a named degraded state; the resolve action is pre-emptively disabled with an explanation rather than allowed to fail mid-submission	Work other queue items; the action re-enables automatically when PVQ recovers, with no reload and no re-authentication
Validation failure on the resolution form	Error summary at the top of the form with in-page links to the offending fields; the narrative is preserved verbatim	Fix the named field and resubmit
Session timeout warning mid-narrative	SCR-06 modal with a live countdown and "Stay signed in"; entered data is not discarded	Extend in place; if the session does end, re-authentication returns him to the requested item via <code>returnTo</code>
ISS-2207 already resolved (demo re-run without reset)	The issue detail shows its resolved state and disposition; the resolve action is absent because the server computes it as unavailable	Run the reset command, or drive the journey on the second seeded open issue (precondition P7)
Issue references a case that does not exist (the one intentional seeded orphan)	A "couldn't be confirmed" relationship state naming what could not be resolved — not a broken link and not a stack trace	Documented seed condition; exists so the mismatch handling is demonstrable rather than theoretical
Marcus opens a case assigned to another investigator by direct URL	SCR-30 access denied, non-enumerable, with a correlation ID and exits to his dashboard and queue; the denial is audited	Return to his own queue

Accessibility Notes

*This journey must be completable **keyboard-only and with a screen reader**, start to finish (SM-08). It is the accessibility pass that matters most, because it is the path a reviewer will watch.*

- **Queue table (Stage 3):** proper header scope and caption; sortable headers announce sort state; filter and pagination changes announce a result count. Without this the queue is unusable non-visually.
- **The traversal (Stage 7) is the highest-risk accessibility step in the product.** Moving from the eApp case to the related PVQ issue must produce a **descriptive page-title change** and place focus predictably on the new heading. If focus is dropped to the document top or left on a detached element, a screen-reader user loses precisely the continuity this product claims to deliver.

- **Source attribution** (eApp , PVQ , PDT , IM) must be in the **accessible name** of each row and panel, never conveyed by a colored badge alone (NFR-02).
- **Resolution form (Stage 8):** programmatically associated labels, described-by hint text on the disposition options, required-field indication that is not color-only, inline errors, and an **error summary with focus management and in-page links**. The narrative is long free text; losing it to a validation failure is a trust-destroying event.
- **Dual-system confirmation (Stage 9):** announced via aria-live="polite" with focus moved deliberately to the alert, since this is a result the user is waiting for. Contrast with degraded warnings, which must **never** steal focus.
- **Status vocabulary** — open, resolved, overdue, blocked — distinguishable in grayscale by text and/or icon.
- Usable at **200% zoom** and in a **half-width browser window at 320px equivalent** with no horizontal scrolling, because Marcus routinely runs the browser beside a notes document (NFR-16).

Success Outcome

Marcus completes eApp case → related PVQ issue → resolution → dual-system confirmation in **under three minutes, with exactly one authentication event in the audit log, zero manual identifier re-entry, and both spokes independently returning the updated state through their own APIs** — the success measure of JTBD-01.1 and the primary metric set SM-01 through SM-04, with SM-20 satisfied by the correlated chain.

Feature Touchpoints

Stage	Features
1. Sign in	F0, F1, F13
2. Orient	F4, F15, F2, F17
3. Enter the work	F5, F8, F14, F17
4. Open the case	F2, F6, F8, F9
5. Read the answer	F6, F9, F14
6. Discover the relationship	F7 , F6, F8, F9
7. Traverse	F7 , F1, F3, F13, F14
8. Resolve	F7 , F2, F6, F14
9. See both systems answer	F7 , F10, F13, F16
10. Verify on the case	F7 , F6, F9
11. Leave the record	F13, F6, F10

JRN-01.02: Monday-Morning Triage Across the Whole Caseload

Persona: PER-01 (Marcus Vale) **Scenario:** 0705 on a Monday, laptop on a hotel desk over a VPN. Forty-one open cases, a weekend of accumulation, and a day that will be interrupted a dozen times. Marcus needs one trustworthy answer to "what is assigned to me and what is due first" before he starts making phone calls — and he needs that answer to survive the interruptions, so that when he comes back at 0940 after a source calls him, the queue is still where he left it. This journey is also where a newly registered sixth application (CVS) becomes visible to him without signing out, which is why it is paired with JRN-04.01 in the demo. **Related Jobs:** JTBD-01.2 (*primary*), JTBD-01.3 (*recent activity and the record*), JTBD-04.1 (*as the observing half of the extensibility proof*) **Entry Trigger:** Start of the working day, and again after every interruption of 10–60 minutes. **Demo Path:** Supporting — **Segment 3b.** Marcus's queue is the window an evaluator watches while Priya registers CVS in the other window.

Preconditions (Seed Data Required)

#	Namespace	Required seed state
P1	hub	PER-01's queue contains 28–34 items spanning at least four of five spokes, so pagination at 25 per page is exercised and SM-14 is observable.
P2	im	Overdue assignments owned by PER-01, so the overdue alert rule and due-date sort have real material rather than a contrived single row.
P3	eapp / pvq / pdt	Items owned by PER-01 in each namespace with varied status and priority, so every filter facet returns at least one result — a filter that always returns nothing looks broken.
P4	hub	Default saved view for the investigator role: "assigned to me, due date ascending", applied without configuration.
P5	hub	Two active announcements targeted at the investigator role, one of which is dismissible per user, plus ~4 alerts of distinct types (overdue, newly raised issue, blocked case, approaching due).
P6	cvs	Nine CVS alerts, four of them assignable to PER-01, seeded but unreachable until CVS is registered — the visible consequence of JRN-04.01.

Journey Stages

Stage	Action	Touchpoint	Thinking	Feeling	Pain Point	Opportunity
1. Sign in	Authenticates once by CAC/PIV over a VPN that is having a mediocre morning	SCR-01 → SCR-02	"Please just let me in once."	Impatient	Today: three logins with three session lifetimes; the one he needs third has already expired	One session, honored by every adapter for its whole life (F1)
2. Take the temperature	Reads assigned-work counts, alerts, and due-date widgets	SCR-09 investigator dashboard	"Forty-one open. Four overdue. One new issue. Okay."	Braced but oriented	Today: the shape of his own workload has to be reconstructed from three systems and memory	One composed answer to "what is mine, what is urgent, what changed" (F4)
3. Open the queue	Activates "View all work" and lands on the default view	SCR-13 unified work queue, "assigned to me, due date ascending"	"It already sorted it the way I would have."	Mildly surprised	Today: three lists, three sort orders, none prioritized against the others, so he prioritizes by whoever called him last	A default view that answers the question without configuration; next action ≤ 2 clicks from sign-in

Stage	Action	Touchpoint	Thinking	Feeling	Pain Point	Opportunity
4. Narrow	Adds a source-system filter and a due-date range; chips appear with a clear-all	SCR-13 → filter chips + result count	"Show me just what's overdue this week."	In control	Today: narrowing means re-querying each system separately and comparing by hand	Server-scoped filtering with removable chips and an announced result count (F5, F14)
5. Read a row and commit	Scans source badges and status text, opens the highest-consequence item	SCR-13 row → SCR-14/SCR-15 detail	"PVQ badge, due Thursday. That's the one that blocks a case."	Decisive	Today: priority and due-date semantics differ invisibly between systems, so a sorted list can be wrong in a way he cannot see	Normalized date and priority semantics with visible source attribution on every row
6. Get interrupted, come back	A source calls; he abandons the screen for forty minutes, then returns via the browser	SCR-13 restored with prior filters, sort, and page	"It's exactly where I left it. I don't have to rebuild this."	Relieved	Today: he rebuilds filters and scroll position from scratch a dozen times a day	Queue context preserved across return navigation and interruption (F6)
7. Notice something new arrive	A sixth source appears in the filter list and four new items appear in the queue, correctly attributed — without him signing out	SCR-13 → source filter + new CVS-badged rows	"Where did those come from? Nobody restarted anything."	Curious, then impressed	Today: a new application means a new login, a new URL, and an engineering project that took months	Registry-driven fan-out: a newly registered application reaches an open session on the next poll (F8, F12, SM-12)

System Response by Stage

Stage	System Response	Features
1. Sign in	One server-side session; entitlements computed server-side and returned to build navigation, so the UI never advertises what he cannot do	F0, F1, F2
2. Take the temperature	Dashboard composed per role from live spoke data with widget-level loading states, so a slow spoke never blanks the page; alert rules computed server-side	F4, F15, F16
3. Open the queue	Single fan-out request across all registered adapters with per-source partial-failure tolerance; normalized work-item model with source attribution retained	F5, F8
4. Narrow	Filtering, sorting, and search evaluated server-side and scoped to what this principal may see — never a client-side filter over an over-fetched list	F5, F2, F14
5. Read a row and commit	Resource-level authorization on open; queue context (filters, sort, page) captured for the return trip	F2, F6
6. Get interrupted, come back	Saved view state restored; session still live, or SCR-06 timeout warning with an accessible path back that does not discard anything	F6, F0
7. Notice something new arrive	Registry version poll (≤ 30 s) delivers updated navigation and fan-out targets to an open session; CVS items normalized and badged like any other source, with no special-casing anywhere in the hub	F8, F12, F5

Emotion curve (1 anxious → 5 confident): Sign in **2** (*Monday, VPN, forty-one cases*) → Dashboard **3** → Queue **4** → Narrow **4** → Commit **4** → Return from interruption **5** (*the context survived*) → New source appears **5**.

Key Moments

- **Decision Point — Stage 5.** He commits his morning to an ordering. If the sort is wrong in a way he cannot explain, the whole day is mis-spent and he goes back to checking each system by hand.
- **Delight Opportunity — Stage 6.** Returning after forty minutes to filters, sort, and page exactly as he left them. For an investigator with a severe interruption profile, this is worth more than any feature that looks better in a screenshot.
- **Risk of Abandonment — Stage 3.** A queue that shows fewer sources than he knows he has work in. He will not report it; he will quietly resume checking each system directly, and the product has lost its only value.
- **Demo Moment — Stage 7.** New work arriving in an already-signed-in session without a restart is the visible consequence that makes JRN-04.01's registration real rather than administrative.

Success Exit Criteria

- The queue renders correctly attributed items from **at least four of five** spokes (SM-14), rising to six after CVS registration.
- The correct next item is reachable from the dashboard in **two clicks or fewer**.
- Every filter facet returns at least one result on the default date range; pagination is exercised (> 1 page).
- Returning from an item restores filters, sort, and page.
- Sort changes announce sort state; filter changes announce a result count.
- Newly registered CVS items appear **without signing out or reloading the application** (SM-12).

Failure and Alternate Paths

Condition	What Marcus sees	Recovery
One spoke unavailable	Remaining sources render fully plus a named, quantified warning — see JRN-01.03	Keep working; recovery is automatic
Genuinely empty filter result	A designed empty state explaining what would appear here and offering "clear filters" — visibly different from a degraded state	Clear or widen the filters
Session expires during a long interruption	SCR-06 warning while still signed in; if it lapses, sign-in returns him to the requested path via <code>returnTo</code>	Re-authenticate; queue view state is restored
An item he may see but may not act on	The item opens; the action set is computed server-side and simply does not offer what he is not entitled to do	No wasted effort on an action he cannot complete
Search returns another investigator's case by identifier guess	Denied at resource level, non-enumerably; the denial is audited	Continue within his own scope

Accessibility Notes

- **Keyboard-only:** the entire journey — filter chips, sortable headers, pagination, row activation, clear-all — is reachable and operable by keyboard with a visible focus order that follows the visual order. No keyboard traps in the filter panel.
- **Screen reader:** the queue is a proper data table with caption and header scope; sortable headers announce current sort state; filter and pagination changes announce the new result count via a live region; source attribution is in each row's accessible name.

- Announcements and alerts arriving asynchronously (including CVS appearing at Stage 7) are delivered through live regions **without stealing focus** — he may be mid-form elsewhere.
- Overdue, blocked, and new states are conveyed by text and/or icon, never color alone (NFR-02).
- Usable at 200% zoom and at 320px-equivalent width with no horizontal scrolling; the queue reflows rather than requiring lateral scroll (NFR-16).

Success Outcome

Marcus's unified queue renders correctly attributed items from at least four of five spokes and he opens the correct next item from the dashboard in **two clicks or fewer** — the success measure of JTBD-01.2 (SM-14, SM-05) — with his working context surviving a 40-minute interruption and a newly onboarded sixth application reaching his open session without a sign-out (SM-12).

Feature Touchpoints

Stage	Features
1. Sign in	F0, F1, F2
2. Take the temperature	F4, F15, F16, F17
3. Open the queue	F5, F8, F17
4. Narrow	F5, F2, F14
5. Read a row and commit	F5, F6, F2
6. Get interrupted, come back	F6, F0, F3
7. Notice something new arrive	F8, F12, F5, F16

JRN-01.03: Keep Working While Investigation Management Is Down

Persona: PER-01 (Marcus Vale) **Scenario:** Mid-morning, Marcus is at an employer site working over a VPN. Unknown to him, Investigation Management has just been forced into an unavailable state — in the demo, deliberately, by Priya from the failure-injection controls (JRN-04.03); in production, by an ordinary outage. He refreshes his queue. Today, the symptom of an unhealthy back-end is an empty list, and an empty list that might be a lie is worse than an error: he cannot distinguish "no items" from "the system couldn't be reached," so he stops trusting every screen. This journey is the product's answer: degradation that is **named, quantified, and survivable**, with the other four sources fully usable and recovery that costs him nothing. **Related Jobs:** JTBD-01.4 (*primary*), JTBD-01.2 (*the queue he is still working*) **Entry Trigger:** A spoke becomes unhealthy while Marcus has an open session — detected by the background health monitor, not by Marcus. **Demo Path:**  **SECONDARY SCRIPTED DEMO — Segment 2a.** Driven from the other window by JRN-04.03. The evaluator watches this window.

Preconditions (Seed Data Required)

#	Namespace	Required seed state
P1	im	IM is the designated outage-demonstration spoke. Its seed data must include items owned by PER-01 — including overdue ones — so that removing it visibly changes counts . A degradation demo against an empty source is abstract and unconvincing.
P2	hub	A populated <code>work_item_counts_cache</code> for PER-01 × IM from a prior successful fan-out, so the warning can quantify the gap ("12 items are not shown") rather than fall back to the unquantified copy.
P3	eapp / pvq / pdt	Enough PER-01-owned items in the remaining spokes that the queue is still substantively useful with IM removed — the claim is "the rest of your work is up to date," and it has to be true on screen.
P4	im	At least one in-flight action target owned by PER-01 in IM, so the pre-emptively disabled action state is demonstrable rather than theoretical.
P5	hub	Failure-injection controls available to the administrator identity (SCR-38), demo-scoped and administrator-only.
P6	hub	Circuit-breaker thresholds and health-probe interval configured tightly enough that the demo does not require an awkward silence — degradation and recovery must both be observable inside one health-check interval.

Journey Stages

Stage	Action	Touchpoint	Thinking	Feeling	Pain Point	Opportunity
1. Refresh the queue	Returns to the queue after a call and the page refreshes	SCR-13 unified work queue	"Hang on — that's fewer rows than it was."	Suspicious	Today: an empty or short list is indistinguishable from an outage, so he cannot trust what he is looking at	Partial results by default: what is available renders, what is missing is stated
2. Read the warning	Reads a prominent site-alert: "Investigation Management is unavailable — 12 items are not shown. The rest of your work is up to date."	SCR-13 → degraded-system banner	"Twelve missing, and I know which system. Fine — that's not my problem to solve."	Annoyed but oriented	Today: he calls the administrator or a colleague to ask whether "it's just me," and waits	Degradation named and quantified , so the gap is bounded rather than unknown (NFR-10)
3. Keep working the rest	Filters to eApp and PVQ and continues triaging	SCR-13 → source filter, remaining four sources fully actionable	"Four out of five still works. I can use the morning."	Productive again	Today: one unhealthy back-end stalls the page or blanks the list, and he stops working entirely	One slow or dead spoke never stalls the aggregate request; the rest stays actionable (SM-16)
4. Hit a blocked action	Opens an item whose action targets IM; the action is visibly disabled with a plain explanation	SCR-14 work-item detail → disabled action + explanatory text	"It's telling me now, not after I typed three paragraphs."	Respected	Today: the action is offered, accepted, and fails mid-submission, taking his typed input with it	Actions against an unavailable spoke are pre-emptively disabled with an explanation , never allowed to fail mid-flight

Stage	Action	Touchpoint	Thinking	Feeling	Pain Point	Opportunity
5. Check whether it is just him	Glances at the header status indicator; the warning is consistent everywhere incomplete data is shown	SCR-08 shell header → SCR-09 dashboard degraded notice	"Same message on the dashboard. It's the system, not me."	Settled	Today: each screen fails differently, so he cannot form a single picture of what is wrong	Consistent, specific degraded treatment wherever incomplete data appears (F16)
6. Recovery arrives	The next successful health check restores IM; the warning clears and the missing items reappear	SCR-13 live-region update — no reload, no re-authentication	"It came back on its own. I didn't have to do anything."	Reassured	Today: recovery means a page reload at best and a fresh login at worst, usually discovered by trial and error	Automatic recovery via half-open circuit probing; the user does nothing (SM-17)
7. Resume with a full picture	Clears the filter and sees all five sources and the restored count	SCR-13 full queue + announced result count	"Back to forty-one. Nothing got lost."	Confident	Today: he re-verifies each source by hand before he trusts the total again	Counts reconcile after recovery, so trust is restored rather than merely assumed

System Response by Stage

Stage	System Response	Features
1. Refresh the queue	Fan-out returns HTTP 200 with partial results and a per-source status list; the failing adapter's error is logged with a correlation ID into the integration error log	F5, F16, F8, F11
2. Read the warning	Degraded banner rendered from per-source status, naming the application and quantifying the omission from the last successful count; where no prior count exists the copy omits the number rather than guessing	F16, F14, NFR-10
3. Keep working the rest	Circuit breaker opens after the configured failure threshold, so the hub stops hammering IM and the remaining adapters answer at full speed	F16, F8, F9
4. Hit a blocked action	Server-computed action list marks IM-targeted actions unavailable with a reason; the client renders them disabled with explanatory text rather than hiding them, so the capability is legible	F6, F2, F16
5. Check whether it is just him	Degraded state is surfaced consistently on dashboard, queue, and detail; no error page and no blank screen appears anywhere in the application	F4, F16, F3
6. Recovery arrives	Half-open probing succeeds, the circuit closes, health flips to healthy, polling delivers restored data to the open session; the warning clears through the same live region that raised it	F16, F1
7. Resume with a full picture	Full fan-out resumes; <code>work_item_counts_cache</code> refreshed; result count announced	F5, F16, F14

Emotion curve (1 anxious → 5 confident): Refresh **2** (*the count is wrong*) → Read warning **3** (*named and bounded — the anxiety converts to information*) → Keep working **4** → Blocked action **4** (*told before, not after*) → Cross-check **4** → Recovery **5** → Resume **5**.

Key Moments

- **Decision Point — Stage 2.** Does he keep working or stop and start phoning people? The named, quantified warning is the entire difference. "Investigation Management is unavailable — 12 items are not shown" converts an outage into a bounded, manageable fact.
- **Risk of Abandonment — Stage 1.** If the short list arrives with no explanation, he learns the screen can lie by omission. That trust loss is permanent: he will re-check every source system by hand forever, and the product's only value is gone.
- **Risk — Stage 4.** A warning that steals focus out of a half-written narrative is itself a failure. Degraded notices are polite live-region updates, never focus grabs.
- **Delight Opportunity — Stage 6.** Recovery with **no reload and no re-authentication** is the observable that separates "we handled the error" from "we designed for it" (SM-17).

Success Exit Criteria

- The queue renders the remaining four sources **fully and actionably** with a **specific named warning** (SM-15, SM-16).
- **No error page, blank screen, or stack trace appears anywhere in the application** during the outage (NFR-09, SM-15).
- Actions targeting the unavailable spoke are disabled with an explanation before submission, not failed after it.
- Restoring the spoke clears the warning and restores data **without user reload or re-authentication** (SM-17).
- The degraded warning is announced to assistive technology **without stealing focus**.
- A correctly attributed entry appears in the integration error log within one health-check interval (pairs with JRN-04.02).

Failure and Alternate Paths

Condition	What Marcus sees	Recovery
Two spokes down simultaneously	Both named in the warning with their own counts; the remaining sources still render	Unchanged behavior; the warning composes rather than collapsing into "something went wrong"
Spoke is slow rather than dead	Widget- and section-level loading states with accessible busy announcements; the adapter times out at its configured bound and degrades to the named warning rather than hanging the page	Page remains interactive throughout
Omitted count unknown (no cached prior count)	"Investigation Management is unavailable — some items are not shown." The number is omitted rather than guessed	Count returns on the next successful fan-out
The spoke returns but with errors	Health reports degraded rather than healthy; the warning persists with the degraded wording; the circuit reopens on repeated failure	Automatic, with the error log recording each failure class
Marcus is mid-form in another item when recovery lands	The live region announces the change; focus is not moved ; his input is untouched	He finishes the form and sees the refreshed queue on return

Accessibility Notes

- **Keyboard-only:** every control remains reachable during degradation. Disabled actions stay **focusable and explained** (or are paired with adjacent explanatory text) so a keyboard user learns *why* rather than finding an inert control.

- **Screen reader:** the degraded warning is delivered through a live region with `aria-live="polite"` and is also present in the page's static structure, so a user who arrives after the announcement still encounters it by navigating headings and landmarks. Announcements **must not** move focus.
- The warning names the application in text — health and degraded states are never conveyed by a red/green dot or color alone (NFR-02).
- Loading and busy states carry accessible busy announcements; a skeleton without semantics is a silent screen.
- The distinction between a **designed empty state** and a **degraded state** must be clear non-visually as well as visually: different wording, different structure, not merely a different icon color.
- Usable at 200% zoom and 320px-equivalent width; the banner must not consume the viewport or push the demo banner out of view (NFR-13, NFR-16).

Success Outcome

With a spoke forced offline from the administrator's failure-injection control, Marcus's queue renders the remaining sources plus a **specific named warning with no error page anywhere in the application**, and restoring the spoke clears the warning **without reload or re-authentication** — the success measure of JTBD-01.4 (SM-15, SM-16, SM-17).

Feature Touchpoints

Stage	Features
1. Refresh the queue	F5, F16, F8
2. Read the warning	F16, F14, F5
3. Keep working the rest	F5, F16, F9
4. Hit a blocked action	F6, F2, F16
5. Check whether it is just him	F4, F3, F16
6. Recovery arrives	F16, F1, F11
7. Resume with a full picture	F5, F16, F14

PER-02: Dana Okonkwo

Adjudicator, Consolidated Adjudication Services. 15–25 cases pending determination, 25–90 minutes each. Reads across eApp, PVQ, PDT, and IM to make one decision that must be defensible to a reviewing authority, an appeal, and an inspector general years later. Low interruption frequency, very high context depth.

JRN-02.01: Assemble the Whole Picture on a Subject and Render a Determination

Persona: PER-02 (Dana Okonkwo) **Scenario:** A subject reaches Dana for determination. The case is the one Marcus worked: eApp case `CASE-A-1042` on subject `SUBJ-00418`, with one PVQ issue that was raised against a Section 13A employment answer and has since been dispositioned. To decide, Dana needs four things — the questionnaire as submitted, every issue raised against it **and how each was disposed**, the PDT designation that set the investigation tier, and the IM case

record showing investigative completeness. Today that is four logins, four searches, and an assembly job she performs by hand on paper, which is the actual work and is unbillable. In the unified layer it is one reading path. This journey also carries the product's clearest live RBAC demonstration: Dana opens **the same PVQ issue work item Marcus resolved** and is offered a demonstrably different, server-computed set of actions. **Related Jobs:** JTBD-02.1 (*primary*), JTBD-02.2 (*the determination and its record*) **Entry Trigger:** A case transitions into her adjudicative queue as investigatively complete. **Demo Path:** □ **SECONDARY SCRIPTED DEMO — Segment 4.** The side-by-side action-set contrast with JRN-01.01 is the sharpest zero-trust moment in the product and should be shown immediately after the flagship, while the evaluator still has `ISS-2207` in mind.

Preconditions (Seed Data Required)

#	Namespace	Required seed state
P1	hub	An ADJUDICATOR identity bound to PER-02, <code>org=DCSA-FIELD-OPS-EAST</code> , <code>tier=T5</code> , with <code>CASE-A-1042</code> 's subject inside her adjudicative scope — and without the investigator's <code>case_assignment</code> , so the action-set difference is attribute-driven, not hard-coded.
P2	pvq	<code>ISS-2207</code> in a resolved state with its disposition and full resolution narrative populated. Dana's determination quality depends on reading the narrative, not merely that an issue closed. (<i>If the demo runs Segment 4 after Segment 1, this state is produced live by Marcus — which is the better story.</i>)
P3	pdt	An <code>APPROVED</code> designation for the subject's position with its resulting investigation tier , since adjudicating against the wrong tier is a finding against her.
P4	im	The IM case record for <code>CASE-A-1042</code> showing investigative completeness and its chain of activity.
P5	hub	15–25 items in her pending-determination queue with varied age , including at least one blocked-awaiting-investigator item, so the aging story (JRN-02.02) has material.
P6	hub	An adjudicator dashboard composition visibly different from the investigator's — different widgets, different counts, different navigation (SM-05).
P7	pvq	A second subject whose issue set is genuinely empty , so "no issues exist" and "PVQ is unreachable" can be shown as two different screens (JRN-02.02).

Journey Stages

Stage	Action	Touchpoint	Thinking	Feeling	Pain Point	Opportunity
1. Sign in and settle	Authenticates once, sees a determination-shaped dashboard	SCR-01/SCR-02 → SCR-10 adjudicator dashboard	"Same product as the investigator, but this is my work, not his."	Composed	Today: four logins with mismatched session lifetimes; PVQ has expired by the time she finishes reading eApp	One session for the entire 45-minute read, no credential prompt at any boundary (F1)
2. Pick the determination	Sorts her pending queue by age and opens the oldest	SCR-13 queue, adjudicator default view	"Oldest first. The clock is measured on me."	Purposeful	Today: she learns a case is aging from a management report, not from her workspace	Aging surfaced in the workspace where the decision is made (F5, F4)
3. Read what was submitted	Reads the eApp questionnaire as submitted, section by section	SCR-15 eApp case view, source-badged eApp	"This is the record as he gave it. Good."	Absorbed	Today: she assembles the four sources herself, on paper or in a side document — the assembly <i>is</i> the work	The record assembles itself; her time goes to judging evidence, not collecting it

Stage	Action	Touchpoint	Thinking	Feeling	Pain Point	Opportunity
4. Follow the issue	Opens the related PVQ issue from the related-items panel — no search, no identifier	SCR-15 related items → SCR-16 PVQ issue detail	"And here's the answer it was raised against, one click away."	Impressed	Today: she reads a resolution narrative in PVQ with no ability to see the questionnaire answer it concerns; the relationship exists in the data and not on her screen	Cross-system relationships resolved inline so traversal replaces four searches (F6, F7)
5. Read the disposition and the narrative	Reads the investigator's disposition and his narrative , with actor and timestamp	SCR-16 → disposition panel + activity history	"He interviewed the employer. That's why the date changed. I can rely on this."	Confident in the record	Today: she sees that an issue existed and closed, with no visibility into the reasoning that closed it	The quality of her determination is only as good as the record; the record is now legible (F6, F13)
6. Notice what she cannot do	Sees that the resolve action Marcus used is not offered to her on this same item	SCR-16 → server-computed action set	"I can read this. I can't resolve it. That's correct — and the system knows it, not just the policy."	Trust in the boundary	Today: she relies on convention rather than enforcement, and is never confident she has seen everything she is entitled to see	Action-level authorization computed server-side per item and per principal, re-authorized at execution (F2)
7. Check the tier	Opens the PDT designation and confirms the investigation tier	SCR-17 PDT designation view, source-badged PDT	"Tier 3. The investigation matches the designation."	Careful	Today: a fourth login and a fourth search against a reference she copied from the third system	Designation and resulting tier reachable in the same reading path (F6)
8. Render the determination	Completes the determination form — decision, rationale, supporting references — and submits	SCR-14 determination form (long, high-consequence, validated)	"This has to stand up on appeal in four years."	Deliberate, weighty	Today: she records the determination in one system and documents its basis in a side note, because no single record ties her decision to the material she saw	Determination and its basis captured together with a durable, correlated audit record (F13)
9. Confirm and move on	Reads a confirmation naming what changed and in which system; returns to her queue with context intact	SCR-14 confirmation → SCR-13 restored	"It says which system, and I can see my own entry in the trail."	Settled, finished	Today: she re-opens systems to confirm her own write, because confirmations name nothing and quantify nothing	Confirmation names system and resulting state; her own action appears in her recent activity (F6, F13)

System Response by Stage

Stage	System Response	Features
1. Sign in and settle	One session; entitlements computed server-side produce an adjudicator-specific navigation set and dashboard composition, visibly different from the investigator's	F0, F1, F2, F4
2. Pick the determination	Queue fan-out scoped to her adjudicative assignment and clearance tier; consistent date semantics normalized across sources so the sort is explainable	F5, F2, F8
3. Read what was submitted	eApp content rendered with source attribution present in the accessible name , correct heading hierarchy and landmarks for heading-based navigation across a dense record	F6, F14

Stage	System Response	Features
4. Follow the issue	Related refs resolved live through the PVQ adapter; breadcrumb carries subject and case context; no credential prompt at the boundary	F6, F7, F1, F3
5. Read the disposition and the narrative	Spoke-native history merged with hub audit records into one chronology showing actor, action, timestamp, and originating system	F6, F13
6. Notice what she cannot do	The server computes the action set for this principal on this item : read actions yes, <code>ISSUE.RESOLVE</code> no. A direct API call to the investigator-only resolve endpoint is denied server-side and the denial is audited	F2 , F10, F13
7. Check the tier	PDT adapter read, source-badged, with programmatic grouping so a non-visual reader can tell where eApp content ends and PDT content begins	F6, F9, F14
8. Render the determination	Action re-authorized at execution time; validation returns an error summary with focus management and in-page links; one audit record written before the success response returns , capturing actor identity and the roles/attributes she held at that moment	F2, F6, F13, F14
9. Confirm and move on	Confirmation names the affected system and resulting state; queue filters, sort, and page restored on return	F6, F5, F13

Emotion curve (1 anxious → 5 confident): Sign in **3** → Pick **3** → Read questionnaire **4** → Follow the issue **5** (*the relationship is already on her screen*) → Read the narrative **5** → See the boundary **5** (*enforcement, not convention*) → Check tier **4** → Render determination **2** (*the heaviest moment in the product — a government decision about a person*) → Confirm **5**.

Key Moments

- **Decision Point — Stage 8.** The determination itself. Everything before it is preparation; everything after it is record. A lost narrative here costs an hour and her trust in the form.
- **Delight Opportunity — Stage 4.** The traversal from questionnaire answer to the issue raised against it, without a search and without an identifier. For Dana this replaces the paper assembly that currently *is* her job.
- **Demo Moment — Stage 6.** Opening the same work item as the investigator and seeing a different, server-computed action set. Show it side by side with JRN-01.01, then show the curl denial and the audited entry.
- **Risk of Abandonment — Stage 3/7.** Any field she cannot attribute to a source system. Unattributed data is, to her, unusable data — and one unattributed panel makes her doubt the rest of the screen.
- **Trust-destroying Risk — Stage 6.** An action offered to her that the server then refuses. A UI that advertises what she cannot do makes her distrust every other control in the product.

Success Exit Criteria

- Every artifact needed for one determination — questionnaire, issue disposition and narrative, designation and tier, case status — is reached **without leaving the unified shell and without re-authenticating** (SM-02, SM-04).
- **Source attribution is present on 100%** of work-item rows and detail panels, in the accessible name and not by badge colour alone.
- Dana and Marcus open the same work item and are presented with **different, server-computed action sets**; a direct API call by Dana to the investigator-only resolve action is denied and audited (SM-18).
- Her determination writes **exactly one audit record before the success response returns** (SM-19).

- Cross-system activity for the case renders as a **single correlated chain**, not four disconnected histories (SM-20).

Failure and Alternate Paths

Condition	What Dana sees	Recovery
PVQ unavailable during the read	A specific named warning stating issue data is unavailable — never an empty issue list that reads as "no issues"	She defers the determination rather than deciding on incomplete data; see JRN-02.02
A case outside her organization or above her clearance tier	SCR-30 access denied, non-enumerable, with a correlation ID; the denial is audited	Return to her queue; nothing about the case's existence is disclosed
Validation failure on the determination form	Error summary at the top with in-page links to the offending fields; the rationale narrative is preserved	Correct the named field and resubmit
She needs something the investigator did not record	She uses "request clarification", which marks the case blocked pending investigator response and surfaces as a work item in his queue	Tracked as blocked until his response arrives (JRN-02.02)
She returns the case for additional investigation	Recorded as a distinct audited action with its own confirmation, moving the case back to the investigator	The case leaves her aging queue with an evidenced reason
Session approaches expiry during a 45-minute read	SCR-06 warning with countdown and "Stay signed in"; nothing entered is discarded	Extend in place

Accessibility Notes

- **Keyboard-only and screen-reader completable end to end.** Long-form reading is her core activity, so **heading hierarchy, landmark regions, and descriptive page titles** are what make a dense multi-source adjudication record navigable — power users navigate by heading, not by scrolling.
- **Programmatic grouping** (fieldset/legend, section landmarks) must make it unambiguous non-visually where eApp content ends and PVQ or PDT content begins. Source attribution belongs in the accessible name.
- Sortable, filterable queue tables announce sort state and result counts. She sorts constantly.
- The determination form is long and high-consequence: **error summary with focus management and in-page links to offending fields is mandatory.** A lost narrative is a lost hour and a lost user.
- Status vocabulary — resolved, open, returned, overdue, blocked — distinguishable in **grayscale** by text and/or icon (NFR-02).
- Usable at **200% zoom** without loss of function; magnification users are common in a reading-intensive workforce (NFR-16).
- The degraded-data warning (alternate path) is announced via a live region **without stealing focus** mid-read.

Success Outcome

Dana reaches every artifact needed for one determination without leaving the unified shell and without re-authenticating, with source attribution on 100% of rows and panels — the success measure of JTBD-02.1 — and her determination is recorded as exactly one audit record naming her and the attributes she held at the time, while the investigator-only action remains unavailable to her and denied server-side if attempted directly (JTBD-02.2, SM-18, SM-19).

Feature Touchpoints

Stage	Features
1. Sign in and settle	F0, F1, F2, F4
2. Pick the determination	F5, F4, F2
3. Read what was submitted	F6, F9, F14, F17
4. Follow the issue	F6, F7 , F1, F3
5. Read the disposition and the narrative	F6, F13
6. Notice what she cannot do	F2 , F6, F10, F13
7. Check the tier	F6, F9, F14
8. Render the determination	F6, F2, F13, F14
9. Confirm and move on	F6, F5, F13

JRN-02.02: Manage the Clock — and Refuse to Decide on Data That Is Quietly Missing

Persona: PER-02 (Dana Okonkwo) **Scenario:** Last week of the month. Dana's timeliness metrics are about to be reported, and she needs to know which of her pending determinations are aging, which are blocked waiting on someone else, and which she can actually close today. Partway through the sweep she opens a subject whose issue list is empty — and this is the moment the product either earns her permanently or loses her permanently. An empty list that is actually an outage would be a career-relevant error, made invisibly. This journey covers both halves of her clock problem: surfacing aging and blocked work without building a report, and being able to tell a genuine absence from an unreachable source. **Related Jobs:** JTBD-02.3 (*aging and blocked*), JTBD-02.4 (*never decide on missing data*) **Entry Trigger:** End-of-month timeliness push; and, within it, opening any subject whose issue list renders empty. **Demo Path:** Supporting — **Segment 2c.** Run immediately after JRN-01.03 while PVQ (or IM) is still injected as unavailable: the same outage that degrades the investigator's queue produces a *different and equally specific* consequence for the adjudicator. One induced failure, two persona-appropriate responses.

Preconditions (Seed Data Required)

#	Namespace	Required seed state
P1	hub	Pending determinations for PER-02 spanning a range of ages , including at least two past their timeliness threshold and two approaching it, so "aging" is visible without contrivance.
P2	hub	At least one determination in a blocked-pending-investigator-response state, created by a seeded clarification request, so "waiting on someone else" is distinguishable from "waiting on me".
P3	pvq	A subject with a genuinely empty issue set. This is the control case: the designed empty state must be demonstrably different from the degraded state, and that comparison is impossible without a true empty.

#	Namespace	Required seed state
P4	hub	Alert rules for overdue, approaching-due, and blocked/stalled computed server-side, each resolving to a real, populated destination — an alert that links to an empty page is worse than no alert.
P5	hub	Failure-injection control able to force PVQ specifically (not only IM) into unavailable, so the adjudicator-facing consequence is demonstrable.
P6	im / eapp / pdt	Remaining sources populated for her scope so the queue stays useful during the injected PVQ outage.

Journey Stages

Stage	Action	Touchpoint	Thinking	Feeling	Pain Point	Opportunity
1. Ask the clock question	Opens her dashboard and reads the aging and blocked widgets	SCR-10 adjudicator dashboard → alerts panel	"What's aging, and what's not actually mine to move?"	Businesslike, slightly braced	Today: aging is invisible until it is a metric someone else reports; she learns a case is late from a report	Aging and blocked computed server-side and surfaced continuously, not retrospectively (F4, F15)
2. Separate mine from theirs	Filters the queue to "blocked" and reads what each is waiting on	SCR-13 queue → status filter + blocked reason	"Three are waiting on investigators. Those aren't my delay."	Relieved of some guilt	Today: she tracks "waiting on investigator" in memory or a side list, because nothing models the wait	Blocked state modelled and attributed, so the clock is managed by evidence (F6, F15)
3. Sort by age	Sorts by age across sources and confirms the ordering is explainable	SCR-13 → sortable headers announcing sort state	"The dates mean the same thing across all of these — I can defend this order."	In control	Today: date semantics differ silently between systems, so a sorted list can be wrong in a way she cannot see	Consistent, normalized date semantics regardless of which spoke supplied the row (F5)
4. Open the next determination	Opens the oldest and begins the read	SCR-15 eApp case view	"Let's close this one before the month ends."	Focused	Today: the read begins with four logins and an assembly job	One reading path (see JRN-02.01)
5. Hit the empty issue list	Reaches the issues section and finds nothing listed	SCR-15 → related issues panel	"No issues at all? On a Tier 3? That's either very clean or very wrong."	Suspicion — the pivotal moment	Today: an empty PVQ list looks identical whether the subject has no issues or the service is down, and the cost of guessing wrong is asymmetric and invisible	The screen must answer the question she is already asking, before she asks it
6. Read which one it is	Reads either a designed empty state ("No issues have been raised against this subject.") or a named degraded warning ("Issue data from PVQ is unavailable — issues for this subject are not shown.")	SCR-15 → empty state or degraded banner (visibly and semantically different)	"It's the system, not the subject. I'm not deciding this today."	Vindicated, then decisive	Today: there is no signal either way, so she either delays everything or accepts the risk silently	Absence and outage are two different screens with two different sentences (F16, NFR-10)

Stage	Action	Touchpoint	Thinking	Feeling	Pain Point	Opportunity
7. Act on the distinction	Defers the determination with a reason, or — on a genuine empty — proceeds and decides	SCR-14 action set (determination actions disabled while the source is unavailable, with explanation)	"Deferring with a recorded reason is defensible. Deciding blind is not."	Professionally safe	Today: she delays cases she is unsure about, adding to the aging problem she is separately measured on	Actions that depend on an unavailable source are disabled with an explanation rather than permitted to fail (F16)

System Response by Stage

Stage	System Response	Features
1. Ask the clock question	Alert rules for overdue, approaching-due, and blocked/stalled evaluated server-side over aggregated spoke data; each alert deep-links to the item that produced it; alert counts reconcile with the queue	F4, F15, F5
2. Separate mine from theirs	Blocked state and its reason carried on the normalized work item; clarification requests tracked until the investigator responds	F6, F15, F8
3. Sort by age	Date fields normalized at the adapter boundary so cross-source sorting is semantically valid; sort state announced	F5, F8, F14
4. Open the next determination	Resource-level authorization; full detail assembled across adapters with source attribution	F2, F6
5. Hit the empty issue list	Per-source status accompanies every aggregated panel — the UI always knows whether a source answered, failed, or answered with nothing	F16, F6, F8
6. Read which one it is	Two distinct, designed states with different copy and different structure: empty vs. degraded. Degraded names the application and what is missing; announced via live region without stealing focus	F16 , F14, NFR-10
7. Act on the distinction	Determination actions that depend on the unavailable source are pre-emptively disabled with a reason; the deferral is a real audited action with its own record	F16, F6, F2, F13

Emotion curve (1 anxious → 5 confident): Dashboard **3** → Separate blocked **4** → Sort by age **4** → Open determination **4** → **Empty issue list 1** (*the sharpest trust dip in the entire document — she is one screen away from a career-relevant error*) → Read which one it is **4** → Act **5**.

Key Moments

- **Decision Point — Stage 6/7.** Decide or defer. The entire value of the degraded-state design is compressed into whether this one screen tells her which world she is in.
- **Trust-destroying Risk — Stage 5.** One instance of an empty list that turns out to have been an outage and the trust loss is **permanent and total** for this persona. She will re-check every source by hand forever, and the product's aggregation value goes to zero.
- **Risk of Abandonment — Stage 1.** An aging view she has to build herself is the workaround she already has; she will keep the spreadsheet and stop opening the dashboard.
- **Delight Opportunity — Stage 2.** Being able to say "three of my late cases are waiting on investigators, and here is the evidence" turns a monthly defensive conversation into a two-minute factual one.

Success Exit Criteria

- Her dashboard and queue answer **"what is aging and what is blocked"** without her constructing a report, and **every alert resolves to a real, populated destination** (SM-05, SM-06).
- With PVQ forced offline, the subject record shows a **specific named warning that issue data is unavailable — and never an empty issue list** (SM-15, NFR-10).
- The designed empty state and the degraded state are distinguishable **both visually and non-visually**, by wording and structure rather than by colour or icon alone.
- Actions depending on an unavailable source are disabled with an explanation before submission.
- A generic "something went wrong" appears nowhere: every incomplete view names the system and the gap.

Failure and Alternate Paths

Condition	What Dana sees	Recovery
PVQ recovers while she is still on the page	The warning clears through the same live region and the issue list populates — no reload, no re-authentication ; focus is not moved	She continues the read with a complete record
Alert count disagrees with the queue	<i>Must not occur.</i> Counts are computed from the same server-side evaluation as the queue rows; a mismatch is a defect, not a display quirk	Treated as a P0 bug — reconciling counts is a hiring criterion for this persona
Blocked item's investigator has responded	The blocked state clears and the item returns to "waiting on me" with the response visible in activity history	She resumes the determination
She defers a determination	Recorded as an audited action with a reason; the item remains in her queue flagged as deferred-pending-source	Revisit when health is restored
Both PVQ and IM unavailable	Both named, each with its own missing-data statement; she is never left inferring which part of the record is incomplete	Defer; no partial-picture determination is offered

Accessibility Notes

- **Keyboard-only and screen-reader completable end to end**, including the sort, the status filter, and the deferral action.
- **The empty-vs-degraded distinction must be legible non-visually.** Different sentences and different structure — not the same container with a different icon colour. A screen-reader user must reach the same conclusion Dana does, in the same second.
- Degraded warnings arrive via `aria-live="polite"` and **must not steal focus** mid-read; they also persist in the static page structure so a user who arrives later encounters them by heading or landmark navigation.
- Aging and blocked status conveyed by text and/or icon, never colour alone — "overdue," "blocked," "approaching due" must be distinguishable in grayscale (NFR-02).
- Queue sort and filter changes announce sort state and result count.
- Disabled determination actions remain discoverable with their explanation programmatically associated, so a keyboard user learns *why* rather than encountering an inert control.
- Usable at 200% zoom without loss of function (NFR-16).

Success Outcome

Dana’s dashboard and queue answer "what is aging and what is blocked" without her building a report, every alert resolves to a real destination (JTBD-02.3, SM-05, SM-06) — and with PVQ forced offline she sees a specific named warning that issue data is unavailable rather than an empty list she could misread as "no issues exist" (JTBD-02.4, SM-15, NFR-10).

Feature Touchpoints

Stage	Features
1. Ask the clock question	F4, F15, F5
2. Separate mine from theirs	F6, F15, F5
3. Sort by age	F5, F8, F14
4. Open the next determination	F6, F2, F9
5. Hit the empty issue list	F16 , F6, F8
6. Read which one it is	F16 , F14, F9
7. Act on the distinction	F16, F6, F2, F13

PER-03: Renée Ashford

The individual being vetted — a private-sector systems engineer at a cleared defense contractor undergoing an initial Tier 3 investigation. Her start date on billable work depends on this process completing. She has no mental model of DCSA's internal system boundaries and must never need one. Short, anxious, frequently interrupted sessions, often on a phone, sometimes eleven days apart.

JRN-03.01: On a Phone at Lunch — Where Do I Stand, What Do I Owe, Do It Now

Persona: PER-03 (Renée Ashford) **Scenario:** It has been eleven days since Renée last looked. She is eating lunch at her desk with her phone in one hand. She wants one thing: a plain-language answer to *where am I in this process and what do I owe you next*. She has forgotten the interface entirely and will not learn it again. She signs in with a non-CAC method — she has no CAC, will never have one, and is not a government employee — reads her status in a sentence she understands, finds one outstanding task with a due date, completes it, and sees the status change as a result. She never learns that the answer was assembled from two systems, and she never sees anything belonging to anyone else. **Related Jobs:** JTBD-03.1 (*where do I stand*), JTBD-03.2 (*what do I owe, do it now*), JTBD-03.3 (*notices where she already looks*) **Entry Trigger:** Anxiety plus elapsed time — roughly 70% of her sessions exist for the status question alone. Secondly, a notice issued to her since her last visit. **Demo Path:** Supporting — **Segment 5a**, immediately before the zero-trust demonstration in JRN-03.02. Show it **on a narrow viewport**; the applicant surface is the one an evaluator is most likely to resize.

Preconditions (Seed Data Required)

#	Namespace	Required seed state
P1	hub	An APPLICANT identity bound to PER-03 with subjectRef = SUBJ-00622 and a non-CAC authentication method (generic MFA with a visible deterministic demo code, or ECA). The applicant persona must never require a CAC — she does not have one, and a CAC-only applicant would be a domain error on screen.
P2	iep	A status record for SUBJ-00622 at a mid-process stage, plus at least one unread notice and one OPEN task with a due date in the INFORMATION_REQUESTED shape — this is the action-required path.
P3	eapp	Her own submitted case for SUBJ-00622 with submission date and current state, so the plain-language narrative is assembled from two systems (IEP + eApp) and the "one process, not two system statuses" claim is real.
P4	iep	Prior completed submissions in her history, so Stage 8 and JRN-03.02 have evidence to show rather than an empty list.
P5	pvq	A PVQ issue raised against one of her answers that she must never see. Its existence in the seed is what makes the access boundary demonstrable; its absence from every applicant screen is the test.
P6	hub	A zero-item applicant identity seeded separately, so every applicant empty state ("You don't have anything to do right now.") is demonstrable without emptying Renée's account.
P7	hub	Plain-language copy mapping for every internal state code that could surface on an applicant screen — no tier codes, no INFORMATION_REQUESTED, no SOI, no system names. The mapping is seed/configuration, not per-screen improvisation.

Journey Stages

Stage	Action	Touchpoint	Thinking	Feeling	Pain Point	Opportunity
1. Get in	Signs in with username and a one-time code on her phone	SCR-01 login → SCR-04 → SCR-05 generic MFA	"Please don't make me reset a password again."	Tense, low expectations	Today: one login for the questionnaire, another for the portal that sends notices; she has reset one of them twice, each reset costing an evening	One credential for the whole process, with the simulated auth labelled honestly (F0)
2. Read where she stands	Reads a one-sentence status and a short progress structure, above the fold, on a 320px screen	SCR-11 applicant dashboard → status widget	"Okay — they have it, it's being reviewed, nothing is wrong."	Relief — the whole point of the session	Today: two screens with different terminology for the same state, and she cannot tell whether they describe the same thing	One plain-language progress narrative assembled from IEP and eApp, presented as one process (F4)
3. Understand, not decode	Reads the status without hitting an acronym or an internal state name	SCR-11 → plain-language copy, no tier codes, no system names	"Nothing here I have to look up. Good."	Respected	Today: "Pending SOI transmittal" is presented as a status; it is not a status, it is a barrier	Zero internal system names, tier codes, or state abbreviations without plain-language explanation

Stage	Action	Touchpoint	Thinking	Feeling	Pain Point	Opportunity
4. See what she owes	Sees one outstanding task with a due date, stated as an obligation not a systems inventory	SCR-11 → outstanding tasks list	"One thing. Due Friday. That I can do at lunch."	Focused, faintly anxious about the date	Today: she learns she owes something when the deadline has nearly passed, via a forwarded email from her security officer	A short, ordered, unambiguous task list with due dates (F4, F5)
5. Read the notice	Opens the unread notice sitting in the same place she checks status	SCR-11 notices region → SCR-21 notifications	"It's here, not in an email I lost."	Reassured	Today: notices arrive by email, sometimes second-hand, and get lost; she cannot tell what she has already handled	Notices in-app, retained, with read/unread state (F15)
6. Do the thing	Activates the task, completes the form, uploads/enters what was asked, submits	SCR-18 IEP task → task action form (USWDS, labelled, validated)	"Am I giving them the right thing? ... Yes, it says exactly what they want."	Concentrating; mildly uncomfortable disclosing personal detail	Today: the task has no path from where she found it; she re-reads an email chain to work out which system it refers to	Every task links directly to the action that discharges it — no dead ends, no "contact your security officer" as a primary path (NFR-14)
7. Get a real confirmation	Reads a confirmation naming what was received and what happens next	SCR-18 confirmation	"Received. Reviewed in about two weeks. That's an answer."	Relief, then done	Today: a bare "Submitted" leaves her unsure anything happened, so she calls her security officer anyway	Confirmation names what was received and what happens next (F6)
8. See it land	Returns to the dashboard; the task is gone, the status line reflects the submission, and the item appears in her history	SCR-11 refreshed → SCR-18 submission history	"The screen changed because of something I did. It's actually connected."	Confident	Today: nothing visibly changes, so she assumes nothing happened and checks again tomorrow	Her action changes her status narrative — the system is legibly responsive (F4, F6)

System Response by Stage

Stage	System Response	Features
1. Get in	Non-CAC path establishes one session; the simulation is labelled on screen so nothing implies real credential validation; no language claiming anything was "verified" or "validated"	F0, F1
2. Read where she stands	Dashboard composes IEP status and eApp case state into one progress narrative server-side; the applicant composition is visibly and substantively different from the three mission dashboards (SM-05)	F4 , F5, F8
3. Understand, not decode	Internal state codes translated through the plain-language mapping before they leave the server; raw state names never reach the applicant surface	F4, F14
4. See what she owes	Outstanding tasks aggregated and ordered by due date, each carrying a working deep link to the action that discharges it	F4, F5, F6
5. Read the notice	Notices rendered with read/unread state and an accessible mark-as-read; new notices announced via live region without stealing focus	F15, F14

Stage	System Response	Features
6. Do the thing	Every read is resource-level entitlement-checked against her subjectRef — route-level checks alone would be a defect; the form validates server-side with an accessible error summary; session-timeout warning offers re-authentication without discarding entered data	F2, F6, F0, F14
7. Get a real confirmation	Write executed through the IEP adapter; one audit record written before the success response returns ; confirmation names the received artifact and the next step	F6, F13
8. See it land	Status recomposed from live spoke data; the task leaves the open list; the submission appears in her self-scoped history	F4, F6, F13

Emotion curve (1 anxious → 5 confident): Get in **2** (*eleven days of silence and a password she half-remembers*) → Read status **4** (*the biggest single emotional jump in this document*) → Understand **4** → See what she owes **3** (*a due date is a small spike of stress*) → Read notice **4** → Do the thing **3** (*disclosing personal detail is uncomfortable*) → Confirmation **5** → See it land **5**.

Key Moments

- **Delight Opportunity — Stage 2.** Answering "where am I" in one sentence, above the fold, on a phone, in under 30 seconds. For 70% of her sessions this **is** the product. Everything else is secondary.
- **Decision Point — Stage 6.** Whether she completes the task now or defers it "until the weekend" — which historically means a missed deadline and a phone call to three other people.
- **Risk of Abandonment — Stage 3.** One unexplained acronym or internal state name and she concludes the system is not for her and calls her security officer instead. Jargon is the failure mode, not layout.
- **Risk of Abandonment — Stage 6.** Losing a partly completed form to a silent timeout. She will not start it a third time in the same week.
- **Boundary Moment — throughout.** She sees her own record and nothing else. A PVQ issue exists against one of her answers (precondition P5) and appears **nowhere** on any applicant screen. That asymmetry is a deliberate access-control boundary, not an omission, and it is the setup for JRN-03.02.

Success Exit Criteria

- She can answer "**where am I in this process**" within **30 seconds of signing in, on a 320px-wide viewport, without scrolling past a fold of jargon** (SM-05, NFR-16).
- The screen contains **zero** internal system names, tier codes, or state abbreviations without plain-language explanation.
- **100% of outstanding tasks link directly to the action that discharges them** (SM-06, NFR-14).
- A submission produces a confirmation naming what was received and what happens next, and then appears in her own history.
- Her dashboard is **visibly and substantively different** from the three mission-role dashboards — same product, different composition (SM-05).
- No applicant screen exposes investigative content, investigator identity, PVQ issue items, PDT designations, or IM case records.

Failure and Alternate Paths

Condition	What Renée sees	Recovery
Session times out mid-form	SCR-06 warning with a countdown and an accessible re-authentication path that does not discard entered data	Re-authenticate and continue; nothing re-typed

Condition	What Renée sees	Recovery
Validation error on a personal-detail field	Inline error plus an error summary with focus management and in-page links; plain-language message, no field codes	Fix the named field; entered content preserved
She has nothing outstanding	Designed empty state: "You don't have anything to do right now. We'll let you know if that changes." — never a blank panel	Nothing owed; the session ends reassured
IEP unavailable	A named plain-language warning that part of her status could not be loaded, without jargon and without a raw system name where one can be avoided	Status returns automatically on recovery; nothing she must do
She tries a URL for an investigator page	SCR-30 access denied, non-enumerable, with exits back to her dashboard; the denial is audited	Returns to her own dashboard
A notice arrives while she is signed in	Announced via live region without stealing focus , with unread state	She reads it when she is ready

Accessibility Notes

*This is the **widest accessibility exposure in the product**. Applicants are the general public and the entire cleared workforce; the population includes screen-reader, keyboard-only, magnification, and cognitive/reading-disability users — and users whose disability is directly relevant to what the vetting process is asking about. A failure here is the one most likely to be noticed externally.*

- **Keyboard-only and screen-reader completable end to end**, including sign-in, task completion, and mark-as-read.
- **Plain language is an accessibility requirement, not a tone preference**. Short sentences, defined terms, no unexplained acronyms, one explicit next action.
- **Mobile is a genuine access path, not a responsive checkbox**: usable at **320px width with no horizontal scroll**, at 200% zoom, with adequate touch targets (NFR-16).
- Forms: programmatically associated labels, described-by hint text for anything ambiguous, required-field indication that is **not colour-only**, inline errors, and an error summary with focus management. She is entering personal detail she may find uncomfortable to disclose; a confusing validation error compounds that.
- **Session timeout** needs a clear warning and an accessible re-authentication path that does not discard entered data — her sessions are short and interrupted, and she *will* be timed out (F0).
- Progress is conveyed by **text and structure**, not by a colour-coded graphic alone (NFR-02).
- Notices and asynchronous updates are delivered via live regions **without stealing focus** (F14, F15).
- The non-dismissible demo banner must not consume the top of a 320px viewport to the point where the status answer falls below the fold (NFR-13 vs. SM-05 — a real layout tension worth testing explicitly).

Success Outcome

Renée answers "where am I and what do I owe next" within 30 seconds of signing in on a phone, with zero unexplained internal terms (JTBD-03.1, SM-05) — and 100% of her outstanding tasks link directly to the action that discharges them, with her submission producing a confirmation naming what was received and what happens next, then appearing in her own history (JTBD-03.2, SM-06, NFR-14).

Feature Touchpoints

Stage	Features
1. Get in	F0, F1, F3
2. Read where she stands	F4 , F5, F9, F17
3. Understand, not decode	F4, F14
4. See what she owes	F4, F5, F6
5. Read the notice	F15, F4, F14
6. Do the thing	F6, F2 , F0, F14
7. Get a real confirmation	F6, F13
8. See it land	F4, F6, F13

JRN-03.02: Prove What I Submitted — and Watch the Boundary Around My Record Hold

Persona: PER-03 (Renée Ashford) **Scenario:** Renée's facility security officer emails to ask whether she sent in the employment documentation two weeks ago. She wants to answer with evidence rather than memory — and while she is in there, she wants the quieter reassurance that the most sensitive information she has ever disclosed is visible **only** to her. This journey has two audiences. For Renée it is a five-minute errand ending in relief. For an evaluator it is the **headline zero-trust demonstration in the product**: an authenticated applicant is the strictest access scope in the system, and the denial is shown live — from the UI and then from curl — with the denial itself written to the audit trail. **Related Jobs:** JTBD-03.4 (*primary*), JTBD-03.1 (*the reassurance that keeps her from calling*) **Entry Trigger:** A third party asks her to prove a prior submission; and, for the demo, the evaluator asking "what stops one applicant from reading another's file?" **Demo Path:** ☐ **SECONDARY SCRIPTED DEMO — Segment 5b.** Runs immediately after JRN-03.01. Drive stages 1–3 in the browser, stages 4–6 with **curl against the hub API**, then show the audited denials in the audit viewer as the administrator. Zero-trust claims that are asserted but not demonstrable are worth nothing here.

Preconditions (Seed Data Required)

#	Namespace	Required seed state
P1	iep / eapp	A completed prior submission for SUBJ-00622 with a real timestamp and a retrievable record, so "prove it" has something to prove.
P2	hub	Self-scoped audit/activity records for her identity covering her own submissions — and nothing else . Her history view is audit scoped to self, not the platform trail.
P3	hub	A second applicant (different subjectRef , e.g. SUBJ-00418) with a known, real resource ID. The denial demo requires a resource that genuinely exists: denying a request for something non-existent proves nothing about enumeration.
P4	pvq / im	Investigative content about Renée that exists and is deliberately unreachable from every applicant surface — a PVQ issue raised against her answer, and an IM assignment naming an investigator.
P5	hub	A published API surface an evaluator can call directly (OpenAPI documentation available), with a consistent non-enumerable error contract and correlation IDs on every response.

#	Namespace	Required seed state
P6	hub	Audit coverage for authorization denials , not just successes, so the denials produced during this demo appear in the viewer within seconds.
P7	hub	An administrator identity available in a second window to display the audit trail immediately after the denials are produced.

Journey Stages

Stage	Action	Touchpoint	Thinking	Feeling	Pain Point	Opportunity
1. Look for proof	Signs in and opens her own submission history	SCR-11 dashboard → SCR-18 submission history	"Did I actually send that, or did I just mean to?"	Uncertain, slightly defensive	Today: she searches personal email to prove what she submitted, and keeps her own copies because she cannot retrieve them from the system	A self-scoped history showing what she sent and when, with links back to the item (F6, F13 self-scoped)
2. Find the entry	Reads the dated entry: what was submitted, when, and its current state	SCR-18 history row → item link	"Sent on the 3rd. There it is, with a date."	Relief	Today: she answers from memory and hopes; her security officer re-asks a week later	Evidence, not recollection — answerable in one screenshot
3. Check the edges	Notices there is nothing on any of her screens about investigators, findings, or issues raised against her answers	SCR-11 / SCR-18 — investigative content absent by design	"This is my file. Just mine. Nothing about anybody else, and nothing I shouldn't have."	Quietly reassured	Today: no visible assurance that the system shows her only her own record, which is unsettling given what the record contains	Privacy legibility: the scope she is in is evident from the shape of what she can reach
4. (Evaluator) Try the side door — wrong subject	An authenticated applicant session calls a legitimate endpoint with another subject's real resource ID	Hub API via curl (GET /api/work-items/{other-subject-item})	(Evaluator:) "Does the check happen on the resource, or only on the route?"	Evaluator: testing	Today: each system interprets scope on its own terms; nobody can demonstrate the boundary, only describe it	Resource-level entitlement check against her subjectRef on every read — route-level checks alone are a defect (F2)
5. (Evaluator) Try the side door — wrong role	The same session calls the investigator-only resolve endpoint directly	Hub API via curl (POST /api/.../actions/RESOLVE_ISSUE)	(Evaluator:) "And if she just calls the mission endpoint?"	Evaluator: testing	Today: authorization lives partly in each UI, so bypassing the UI is an untested path	Server-side authorization on every request; client-side hiding is presentation, never the control (NFR-04, NFR-05)
6. (Evaluator) Read the denials	The administrator opens the audit viewer and filters to the last minute	SCR-33 audit viewer → SCR-34 record detail	(Evaluator:) "So the attempt is on the record, not just blocked."	Evaluator: convinced	Today: denials are not recorded anywhere reconstructable, so an attempted breach leaves no trail	Denials audited with actor, target, outcome, and correlation ID (F13)

System Response by Stage

Stage	System Response	Features
1. Look for proof	Audit/activity query scoped to her identity server-side; she cannot widen it by parameter, because the scope is applied server-side rather than requested by the client	F13, F2, F6
2. Find the entry	Submission record read back through the owning adapter with its real timestamp and current state, linked to the item	F6, F9
3. Check the edges	Navigation and entitlements are server-computed: investigative surfaces are not rendered, not linked, and not reachable by URL. The PVQ issue against her answer exists in the seed and resolves for mission roles only	F2 , F3, F4
4. Wrong subject	Denied server-side with a consistent non-enumerable error that does not reveal whether the resource exists, plus a correlation ID; the denial is written to the audit trail	F2 , F10, F13
5. Wrong role	Denied server-side with the identical error shape and status as the wrong-subject case — the two failures are indistinguishable to the caller, which is the point	F2 , F10, F13
6. Read the denials	Audit viewer returns both denial records with actor identity, roles/attributes held at the time, action type, target system and resource, outcome DENIED , and correlation ID, filterable and paginated	F13, F11

Emotion curve (1 anxious → 5 confident): Look for proof **2** (*she genuinely is not sure*) → Find the entry **5** (*evidence, immediately*) → Check the edges **5** → (evaluator segment — *Renée's journey has ended reassured; stages 4–6 are the reviewer's curve: skeptical 2 → tested 3 → convinced 5*).

Key Moments

- **Delight Opportunity — Stage 2.** Answering her security officer with a dated record in under a minute. Small errand, disproportionate reassurance, and it removes a support contact that today costs three other people time.
- **Demo Moment — Stages 4–6.** This is the product's **headline zero-trust demonstration** (PRD F2 acceptance signal, SM-18). Show it live with curl. An evaluator who has been told about server-side authorization and shown it are two different evaluators.
- **Risk — Stage 4/5.** An error that discloses whether another record exists. A 404-vs-403 distinction, a differing message, or even a differing response time is an enumeration oracle and fails this journey regardless of whether access was granted.
- **Trust-destroying Risk — Stage 3.** Any glimpse of another individual's information, however incidental — a name in a shared header, a count that includes others, an autocomplete leak. One occurrence is unrecoverable for this persona and externally reportable.
- **Risk — Stage 1.** A history that omits something she knows she submitted. A partial record is worse than none, because she will trust it.

Success Exit Criteria

- Her submission history shows **everything she submitted and nothing belonging to anyone else**, retrievable in under a minute.
- An authenticated applicant calling an investigator-only endpoint, **or a legitimate endpoint with another subject's resource ID**, is denied **server-side** with a **consistent non-enumerable error**, and the denial **appears in the audit trail** — demonstrable live via curl (SM-18, PRD F2 acceptance signal).
- The two denial classes are **indistinguishable to the caller** in status, body, and timing.

- No applicant screen or reachable URL exposes investigative content, investigator identity, PVQ issue items, PDT designations, IM case records, or adjudicative deliberation.
- Every denial carries a correlation ID that the administrator can resolve in the audit viewer (JRN-04.02).

Failure and Alternate Paths

Condition	What happens	Recovery
She follows a stale deep link to an item that is no longer hers	SCR-30 access denied, non-enumerable, with exits to her dashboard; denial audited	Returns to her own dashboard
She guesses another subject's ID in the UI	Identical denial to the API path — the UI is not a second authorization surface, it just renders what the server permits	Same
An endpoint is discovered that authorizes only at route level	This is a defect, not an alternate path. Negative-path tests must cover every endpoint with an unauthorized resource identifier, not merely an unauthorized role (NFR-04, F19)	Fix before demo; it invalidates the segment
Audit write fails during a denial	The denial still stands; the audit failure is itself surfaced and logged — an action that cannot be audited does not complete, and that rule applies to recording denials too	Administrator triages via JRN-04.02
Evaluator asks to try it themselves	Published API documentation and a running endpoint are available; nothing about the demo depends on the evaluator not looking	Encouraged — this persona's scope is the one that best survives inspection

Accessibility Notes

- **Keyboard-only and screen-reader completable end to end** for the applicant-facing stages (1–3): history table, row links, and the return path.
- The **history table** is a proper data table with caption, header scope, and an announced result count; dates are readable as text, not as a graphic or a relative-only phrase ("2 weeks ago") that loses precision when precision is the point.
- **Access-denied pages must be accessible and plain-language:** a clear heading, an explanation that does not leak, a selectable correlation ID, and at least two working exits. A denial is a page like any other and is in scope for 508.
- The correlation ID on a denial must be **selectable, copyable text** — she may be asked to quote it, and the administrator certainly will (JRN-04.02).
- No security state is conveyed by colour alone; "you don't have access" is a sentence, not a red box (NFR-02).
- Usable at 320px width and 200% zoom, since this errand will happen on her phone (NFR-16).

Success Outcome

Renée produces evidence of a prior submission in under a minute from a self-scoped history — and an authenticated applicant calling an investigator-only endpoint, or a legitimate endpoint with another subject's resource ID, is denied server-side with a consistent non-enumerable error whose denial appears in the audit trail: the success measure of JTBD-03.4 and the PRD F2 acceptance signal (SM-18).

Feature Touchpoints

Stage	Features
1. Look for proof	F6, F13, F2
2. Find the entry	F6, F9, F17
3. Check the edges	F2 , F3, F4
4. Wrong subject	F2 , F10, F13, F19
5. Wrong role	F2 , F10, F13, F19
6. Read the denials	F13, F11, F10

PER-04: Priya Raghunathan

Platform administrator for the unified layer itself, in the Integrated Enterprise program office. Not a mission user — she is responsible for the thing all three other personas depend on. Expert: reads OpenAPI docs, calls endpoints with curl, and tries to break the authorization model on purpose. She is the reviewer archetype an evaluator will emulate. **Administrators are not exempt** — her own actions are authorized server-side and audited, and platform privilege is explicitly not privilege over mission content.

JRN-04.01: Register the Sixth Application Live, Through the UI

Persona: PER-04 (Priya Raghunathan) **Scenario:** A sixth application — the Continuous Vetting Service — is running, conformant, and completely invisible to the platform: no registry row, nothing in navigation, nothing in any queue, nothing in the console. Priya onboards it herself, in the UI, in one sitting, while an investigator stays signed in at the next desk. Identity, connection, a live connection test, capabilities auto-discovered from the application's own `describe()`, role access, review, submit. No pull request, no deployment, no restart. Ninety seconds later CVS is in the inventory, being health-probed, in role navigation, and its work items are in Marcus's queue — and he never signed out. This is the product's **extensibility proof**, and it is the claim that determines whether the platform's premise holds beyond the five systems it launched with. **Related Jobs:** JTBD-04.1 (*primary*), JTBD-04.2 (*it starts reporting health*), JTBD-04.4 (*the registration is audited*) **Entry Trigger:** A new mission application becomes available to join the platform — in the demo, the moment an evaluator asks "what does it take to add another one?" **Demo Path:** □ **SECONDARY SCRIPTED DEMO — Segment 3a.** Run **two windows side by side:** Priya registering in one, Marcus's already-signed-in queue in the other (JRN-01.02, Stage 7). The consequence in the second window is the proof; the wizard alone is just a form.

Preconditions (Seed Data Required)

#	Namespace	Required seed state
P1	hub	An ADMINISTRATOR identity bound to PER-04, <code>org=DCSA-HQ</code> , with console, registry, audit, and failure-injection entitlements — and no access to mission work-item content.

#	Namespace	Required seed state
P2	hub	Exactly five registered applications at baseline (eApp, IEP, PVQ, PDT, IM). CVS is deliberately absent from the registry . Its absence from navigation, queue, and console before registration is itself part of the demonstration.
P3	cvs	The Continuous Vetting Service running as a sixth spoke process with its own data namespace, implementing the full adapter contract and passing the conformance suite — including a working <code>describe()</code> and health endpoint.
P4	cvs	Nine seeded CVS_ALERT work items, four assignable to the investigator persona , so registration produces an immediately visible change in his queue. The demo needs a visible consequence, not just a new console row.
P5	hub	A retired-identifier table and duplicate-identifier validation, so the "identifier already in use" error path is demonstrable without breaking the happy path.
P6	hub	Reset (FR-F17-11) returns CVS to unregistered , so the registration demo is repeatable identically (SM-22).
P7	hub	An investigator session already open at the start of the segment, so "without signing out" is observed rather than claimed.

Journey Stages

Stage	Action	Touchpoint	Thinking	Feeling	Pain Point	Opportunity
1. Show the "before"	Opens the connected-applications inventory: five rows, no CVS anywhere	SCR-22 admin console — connected applications	"Five. Nothing up my sleeve."	Matter-of-fact	Today: there is no inventory at all; what is connected lives in an informal list that drifts from reality	A registry-backed inventory that is the single answer to "what is attached right now" (F11, F8)
2. Start the flow	Activates "Register application" and enters identity: display name, identifier, description, icon token	SCR-28 registration wizard, step 1 of 5	"This is a form, not a ticket to an engineering team."	Purposeful	Today: onboarding is a bespoke project — requirements, custom code, a release, a regression cycle; weeks to months per application	Onboarding as a configuration action she performs herself (F12)
3. Describe the connection	Enters base endpoint, health endpoint, adapter type; accepts pre-filled timeout, retry, backoff, and circuit-breaker defaults	SCR-28 step 2 — connection	"Sensible defaults. I don't have to invent a retry policy to get started."	Comfortable	Today: every integration re-invents timeout and failure policy, differently, in code	Per-adapter resilience policy as registry configuration, not code (F8)
4. Test it before committing	Runs the live connection test ; reads pass/fail with latency, announced in a live region	SCR-28 step 3 — test connection	"It answered. I'm not submitting into the dark."	Increasingly confident	Today: she cannot try an integration before committing to it; an unreachable endpoint is discovered <i>after</i> the work is done	A live test that gates submission — unreachable endpoints are caught here, not later (F12)

Stage	Action	Touchpoint	Thinking	Feeling	Pain Point	Opportunity
5. Let it describe itself	Reviews work-item types and supported actions pre-populated from the application's own describe() , confirms and edits	SCR-28 step 4 — capabilities (auto-discovered)	"It told the platform what it can do. I just confirmed it."	Impressed — the architectural point lands here	Today: capabilities are discovered by reading someone else's documentation, or by asking their team	Capability negotiation via <code>describe()</code> : applications declare themselves (F8, F12)
6. Decide who sees it	Selects which roles may see CVS; nothing pre-selected	SCR-28 step 5 — access	"Investigators and adjudicators. Not applicants."	Deliberate	Today: role visibility for a new system is configured separately in that system, on its own terms	Role-scoped visibility set once, at the hub, and honored everywhere (F2, F8)
7. Review and submit	Reads the full read-only review with per-step edit links, submits	SCR-28 review → confirmation	"Two minutes. No deployment."	Satisfied	Today: the equivalent milestone is a release note, weeks later	Registry row created; <code>registryVersion</code> incremented; the change is audited naming her and the configuration (F12, F13)
8. Watch it come alive	Sees CVS in the inventory and in health monitoring, with a first probe issued immediately rather than at the next interval	SCR-22 inventory → SCR-24 system health	"It's already being monitored. Nobody restarted anything."	Delight	Today: monitoring for a new system is a separate onboarding with a separate team and its own definition of "up"	Registry-driven behavior everywhere: navigation, fan-out, console, and health all read from the same source (F8, F11, F16)
9. Prove the consequence	Looks at the investigator's window: CVS-badged items in his queue, on a session that never ended	(<i>second window</i>) SCR-13 unified work queue	"That's the whole claim, right there."	Vindicated	Today: a new application means a new URL, a new login, and a user who has to be told it exists	An open session receives updated navigation and fan-out on the next registry poll — zero code changes, zero restarts (SM-11, SM-12)

System Response by Stage

Stage	System Response	Features
1. Show the "before"	Inventory rendered from the registry: display name, identifier, adapter type, endpoint, work-item types, supported actions, enabled state, registration date. There is no hard-coded list of applications anywhere in the hub	F11, F8
2. Start the flow	Multi-step form with per-step validation and state preserved across steps; duplicate and retired identifiers rejected with clear guidance	F12, F14
3. Describe the connection	Timeout, retry, backoff, and circuit-breaker fields pre-filled with defaults and stored as registry configuration for this adapter	F12, F8, F16
4. Test it before committing	Hub calls the candidate's health and <code>describe()</code> endpoints live; result displayed and announced via a live region ; the step cannot be passed on failure without an explicit acknowledgement of a warning-level result	F12 , F14, F16
5. Let it describe itself	<code>describe()</code> output parsed into work-item types and supported actions, pre-populated for confirmation; an application supporting fewer actions degrades its UI affordances rather than erroring	F8 , F12

Stage	System Response	Features
6. Decide who sees it	<code>visibleToRoles</code> written to the registry and enforced server-side in entitlements, navigation, and queue fan-out	F2, F8
7. Review and submit	Registry row created, <code>registryVersion</code> incremented, APPLICATION_REGISTERED audit record written naming the administrator, the application ID, and a configuration summary	F12, F13
8. Watch it come alive	Appears at once in inventory, health monitoring (first probe issued immediately), role-scoped navigation, work-queue fan-out, search fan-out, and related-item resolution targets — with no hub restart	F11, F16, F8, F5
9. Prove the consequence	Open investigator sessions receive updated navigation on their next <code>registryVersion</code> poll (≤ 30 s) without signing out or reloading; CVS items are normalized and badged like any other source, with no special-casing anywhere in the hub	F5, F8, F1, F12

Emotion curve (1 anxious → 5 confident): Show the before **4** (*she is in control; this is her console*) → Step 1 **4** → Connection **3** (*endpoints and policy are where mistakes happen*) → **Live test 5** (*it answered — the relief moment*) → Auto-discovery **5** → Access **4** → Submit **4** (*brief: did it take?*) → Inventory and health **5** → Investigator's queue **5** (*peak*).

Key Moments

- **Decision Point — Stage 4.** The live connection test is the gate. Passing it converts "I hope this endpoint is right" into "the platform just talked to it," and it is the step that makes submitting feel safe.
- **Delight Opportunity — Stage 5.** Capability auto-discovery. The application declaring its own work-item types and actions is the moment the architectural claim becomes concrete for a technical evaluator.
- **Demo Moment — Stage 9.** Work items appearing in an already-signed-in investigator's queue. Do not narrate this; switch windows and let it be seen. **Under five minutes, zero code changes, zero restarts** (SM-11, SM-12).
- **Risk of Abandonment — any stage.** Any step that requires a developer, a pull request, a redeploy, or a restart. That is precisely the status quo she is trying to escape, and one such step invalidates the whole journey.
- **Credibility Risk — the whole journey.** A registration flow that works only for the one application shipped with the demo is hard-coded extensibility. The conformance suite and the absence of CVS-specific code in the hub are what make this honest.

Success Exit Criteria

- CVS is registered **live through the UI in under five minutes, with zero code changes and zero restarts** (SM-11, NFR-11).
- It appears **immediately** in admin inventory, health monitoring, role-scoped navigation, and the unified work queue (SM-12).
- An investigator with an **open session** sees CVS work items **without signing out or reloading** (SM-12).
- The registration writes **one audit record** naming the administrator, the application, and the configuration.
- **De-registration is equally clean:** removing CVS from the registry removes it from navigation, queue, and console with no code change and no errors, and is audited (PRD F8 acceptance signal).
- Reset returns the system to the unregistered state and the demo repeats identically (SM-22).

Failure and Alternate Paths

Condition	What Priya sees	Recovery
Endpoint unreachable at the test step	A named failure with latency and error class, and the step does not pass — submission is blocked	Correct the endpoint and re-test; nothing was written
<code>describe()</code> unsupported or partial	Capabilities step falls back to manual entry; the application registers with fewer declared actions and its UI affordances degrade accordingly rather than erroring	Register with what is declared; refresh capabilities later via "Test connection"
Duplicate or retired identifier	Validation error with clear guidance; identifiers are never reused because existing work-item IDs and audit records reference them	Choose a new identifier
Step validation failure	Error summary with in-page links; state preserved across steps , so nothing already entered is lost	Fix and continue
Registration succeeds but the spoke then fails	It appears in inventory as degraded/unavailable with an integration-error entry within one health-check interval — registration and health are independent concerns	Triage via JRN-04.02; disable if required
She tries to open a CVS work item's content	Denied — administrative privilege over the platform is not privilege over mission content, and the attempt is itself audited	Console-level operation only; this separation is a defensible-design point worth stating aloud in the demo

Accessibility Notes

An operations console is a frequent and unexamined accessibility failure point precisely because it is assumed to be "internal tooling." It is not exempt from Section 508.

- **Keyboard-only and screen-reader completable end to end**, including the wizard, the live test, and the confirmation.
- The multi-step wizard needs **accessible step indication** (current step programmatically conveyed, not just visually), per-step validation with an error summary and focus management, and preserved state across steps.
- The **live connection test must be keyboard-operable and its result announced via a live region** — a result that only appears visually leaves a non-visual user unable to know whether the gate has opened.
- No reliance on drag, hover, or pointer-only interaction anywhere in the flow, including the icon-token picker.
- **Health status is never colour-only**: healthy / degraded / unavailable each need a text label and/or icon. A red/green dot alone is both a conformance failure and an operational hazard (NFR-02).
- Console data tables — inventory, health history — follow the same accessible patterns as the work queue: header scope, captions, sort-state announcement, accessible pagination, announced result counts.
- **Destructive actions** (disable, de-register) use accessible confirmation dialogs that trap focus correctly and restore it on close.
- Endpoints, identifiers, and correlation IDs are **selectable, copyable text**, never truncated cells without an accessible full value.

Success Outcome

Priya registers the sixth application live through the UI **in under five minutes with zero code changes and zero restarts**, and the investigator sees its work items in his unified queue **without signing out** — the success measure of JTBD-04.1 (SM-11, SM-12, NFR-11) and the PRD F12 acceptance signal.

Feature Touchpoints

Stage	Features
1. Show the "before"	F11, F8
2. Start the flow	F12 , F14
3. Describe the connection	F12 , F8, F16
4. Test it before committing	F12 , F14, F16, F9
5. Let it describe itself	F8 , F12, F9
6. Decide who sees it	F2, F8, F12
7. Review and submit	F12 , F13, F14
8. Watch it come alive	F11, F16, F8
9. Prove the consequence	F5, F8, F1, F12

JRN-04.02: Triage a Reported Integration Failure — Health View → Error Log → Correlation ID → Audit Chain

Persona: PER-04 (Priya Raghunathan) **Scenario:** 0916. An investigator messages: "my queue is missing things and an action wouldn't submit — here's the reference number off the error screen." Priya has minutes, not an afternoon, to establish which of four possibilities this is: a spoke outage, an adapter misconfiguration, an authorization denial, or user error. Today that means exporting logs from the hub and every candidate system and lining up timestamps by hand, with nothing tying the records to one another — because nothing ever issued a correlation ID. In the unified layer she starts at the health view, confirms the symptom in the integration-issue log, takes the correlation ID the user already gave her, and follows it into the audit chain where the whole cross-system action reads as one narrative. Then she decides whether to contain. **Related Jobs:** JTBD-04.2 (*know before anyone complains*), JTBD-04.3 (*classify in minutes from one place*), JTBD-04.4 (*account for who did what, to what, when*) **Entry Trigger:** A user report carrying a correlation ID — or, better, the start-of-day health check that catches the same failure before the report arrives. **Demo Path:** Supporting — **Segment 6**, the accountability closer. Runs best immediately after Segment 2 (resilience), because the injected outage has already produced real error-log entries and a real degraded history to triage. Also the natural place to show the denials produced in JRN-03.02.

Preconditions (Seed Data Required)

#	Namespace	Required seed state
P1	hub	Health history for each registered application with at least one prior degraded episode , so the health view shows a history rather than a single green row — a monitor with no past is not credible.
P2	hub	Pre-seeded integration error entries across more than one application and more than one error class (timeout, connection refused, contract violation), so filtering is meaningful rather than decorative.
P3	hub	~120 pre-seeded audit events so the viewer, its filters, and its pagination are populated on first load.

#	Namespace	Required seed state
P4	hub	At least one complete correlated chain in the seed — a multi-spoke action under a single correlation ID — so the chain view is demonstrable even before the live demo produces one.
P5	hub	Audit coverage for authorization denials , adapter failures, authentication events, and registration changes — not only successful mutations.
P6	hub	The administrator's own prior console actions present in the trail, attributed to her — "administrators are not exempt" must be visible, not just stated.
P7	hub	Correlation IDs surfaced in user-facing error states , so the premise of this journey (a user hands her an ID) is real.
P8	hub	A disable/enable capability on a registered application, audited, with immediate effect on user-facing navigation and queue.

Journey Stages

Stage	Action	Touchpoint	Thinking	Feeling	Pain Point	Opportunity
1. Start where the truth is	Opens the health view and reads status, last successful check, latency, and recent history for all applications	SCR-24 admin console — system health	"Is this one system or the whole platform?"	Alert, methodical	Today: health is checked per system, by different teams, with different definitions of "up," and the first signal is a phone call	One consolidated health surface, read from the registry, refreshed by a background monitor (F11, F16)
2. Confirm the symptom	Sees the implicated application degraded with a recent transition time	SCR-24 → application row + check history	"Degraded at 09:07. His message was 09:14. That fits."	Focused, quietly relieved it is explainable	Today: she cannot tell whether the user's complaint and the system's state are the same event	Time-ordered health history makes symptom and cause line up without guessing
3. Read the failures	Filters the integration-issue log by application and time window	SCR-25 admin console — integration issues	"Connection refused, twelve times, one operation. That's the spoke, not the adapter."	Confident in the classification	Today: diagnosis is manual log correlation across the hub and every candidate system, by hand, in a spreadsheet	Filterable error log with timestamp, application, operation, error class, correlation ID, and affected principal (F11)
4. Take the thread	Pastes the correlation ID the user gave her and filters on it	SCR-25 filter → SCR-33 audit viewer	"One identifier. Let's see how far it actually goes."	Testing the claim	Today: there is no correlation ID because nothing ever issued one; the trace stops wherever she stops assembling it	Correlation IDs propagated end to end, browser → hub → spoke, into both audit and error logs (NFR-15)
5. Read it as one story	Opens the chain view and reads every hub and adapter operation belonging to that one user action	SCR-34 audit record detail & chain view	"Read the case, read the issue, wrote PVQ, failed on eApp. That's the whole thing, in order."	Satisfaction — the answer is a record, not a reconstruction	Today: the result is an assertion she has to ask people to believe, heavily caveated	A cross-system action retrievable as a single correlated chain (SM-20)

Stage	Action	Touchpoint	Thinking	Feeling	Pain Point	Opportunity
6. Classify and answer	Determines it is a spoke outage, not misconfiguration, not denial, not user error — and says so, with references	SCR-34 → copy correlation ID into the ticket	"Spoke outage. Twelve items withheld. He wasn't doing anything wrong."	Decisive	Today: she guesses the failure class from the symptom and escalates to whichever team seems most likely	Four candidate causes distinguished from one place , in minutes
7. Contain if needed	Disables the implicated application from the console, with an accessible confirmation	SCR-23 application detail → disable (confirmed)	"Take it out of their queue until it's stable. Better a named absence than intermittent failure."	In control	Today: containment requires someone else's deployment, so the failing integration stays in users' faces	Enable/disable at runtime, audited , and immediately reflected in user navigation and queue (F11, F8)
8. Close the loop	Confirms her own console action appears in the audit trail attributed to her; posts an announcement targeted at the affected roles	SCR-33 audit viewer (filtered to her) → SCR-29 announcements	"My own action is in there under my name. As it should be."	Accountable, done	Today: auditability is per-system, administrative actions least of all, and she communicates by emailing everyone	Administrators are not exempt; targeted announcements replace the email blast (F11, F13, F15)

System Response by Stage

Stage	System Response	Features
1. Start where the truth is	Background health monitor's recorded status, latency, and check history per registered application, with status conveyed by text and/or icon , never a dot alone	F11, F16, F14
2. Confirm the symptom	Check history with transition times; circuit-breaker state visible in its effects — the hub stopped hammering the failing spoke	F16, F8
3. Read the failures	Chronological, filterable adapter-failure log with timestamp, application, operation, error class, correlation ID, and affected principal, linked to the corresponding audit entries	F11 , F13
4. Take the thread	The same correlation ID resolves in both the error log and the audit trail; it is present as selectable, copyable text , never a truncated cell	F13, F10, NFR-15
5. Read it as one story	Chain view returns every record sharing the correlation ID: hub operations and adapter calls, in sequence, with actor, roles/attributes at the time, target system and resource, and outcome	F13
6. Classify and answer	Denials, adapter failures, authentication events, and successful mutations are all in the same trail, so the four candidate causes are distinguishable without leaving the console	F13, F11, F2
7. Contain if needed	Disable writes an audit record and takes effect immediately in entitlements, navigation, and queue fan-out — no restart; users see a named absence rather than intermittent failure	F11, F8, F13
8. Close the loop	Her console actions are authorized server-side and audited under her identity; announcements are authored with severity, target roles, and effective/expiry dates, rendered on the right dashboards and dismissible per user, never obscuring the demo banner	F11, F13, F15

Emotion curve (1 anxious → 5 confident): Report arrives **2** (*unknown scope, a user waiting, the clock running*) → Health view **3** → Symptom confirmed **4** → Error log **4** → Correlation ID **4** → **Chain view 5** (*the whole action in one place — the moment the afternoon becomes four minutes*) → Classify **5** → Contain **5** → Close the loop **5**.

Key Moments

- **Decision Point — Stage 7.** Contain or wait. Disabling an application is user-visible and audited; she needs to make it quickly, from evidence, and be able to justify it later.
- **Delight Opportunity — Stage 5.** One identifier resolving the entire cross-system chain. This is the difference between an answer that is a record and an answer that is an assertion she asks people to believe.
- **Demo Moment — Stage 8.** Her own console action, in the audit trail, under her name — paired with the fact that she **cannot** read the mission content of the work item she just triaged. "Administrators are not exempt" and "platform privilege is not content privilege" are demonstrable claims, not policy statements.
- **Risk of Abandonment — Stage 4.** A trace that stops at the hub boundary and never reaches the adapter call. If the chain is partial, she goes back to manual log merging and the console becomes a dashboard she does not open.
- **Credibility Risk — Stage 1.** A health view that reports "up" while users are experiencing failures. A dashboard that lies is worse than no dashboard, and this persona will catch it.

Success Exit Criteria

- Given a correlation ID taken from a user-facing error state, she retrieves **the complete chain of hub and adapter operations for that action in a single filtered query** (SM-20).
- An induced adapter failure produces a **correctly attributed entry in the integration error log and a degraded/unavailable health state within one health-check interval** (PRD F11 acceptance signal).
- Disabling the implicated application is reflected in user navigation and queue **immediately** and is **itself audited**.
- Her own console actions appear in the audit trail **attributed to her**.
- An attempt by her to read mission work-item content, or to mutate an audit record through any application path, is **denied and itself audited** — and **no application path to modify or delete an audit record exists for anyone**, verified by test (NFR-07).
- Total time from report to classification: **minutes**, from one console.

Failure and Alternate Paths

Condition	What Priya sees	Recovery
The user has no correlation ID	She filters the audit trail by actor, action type, target system, and time window instead — every user-facing error state carries an ID, so this should be rare	Same chain reached by a longer route
The failure was an authorization denial , not an outage	The trail shows outcome DENIED with the principal, the target resource, and the attributes held at the time — user error or a policy question, not an incident	Answer the user; no containment needed
The failure was a contract violation by the spoke	Error class in the log distinguishes it from a timeout or refusal; the adapter conformance suite is the follow-up	Escalate to the application's team with evidence, not a symptom

Condition	What Priya sees	Recovery
Health says healthy but users report failures	Manual "test connection" from the application detail view performs a live check on demand, settling the disagreement immediately	If the live test fails, the monitor's interval or thresholds are the defect
She disables an application by mistake	Re-enable is equally immediate and equally audited; probing restarts at once rather than waiting for the next interval	Both actions are on the record, which is the point
Audit query returns a chain with a gap	A defect, not an alternate path — a write that produced no audit record, or a denial never recorded, is a P0 bug (NFR-06)	Fix before demo; the segment's whole claim depends on completeness

Accessibility Notes

- **Keyboard-only and screen-reader completable end to end** — health view, error log, audit viewer, chain view, disable confirmation, and announcement authoring.
- **The audit viewer is the densest table in the product and the one an evaluator is most likely to inspect.** It needs header scope, a caption, sort-state announcement, accessible pagination, and announced result counts — the same patterns as the work queue, not a lesser console variant.
- **Correlation IDs must be selectable, copyable text with an accessible full value** — never images, never truncated cells without the full string available. She pastes them into tickets constantly.
- Health status conveyed by **text label and/or icon**, never colour alone. A red/green dot is a conformance failure and an operational hazard (NFR-02).
- **Destructive actions** (disable, de-register) use confirmation dialogs that **trap focus correctly and restore it on close**.
- Filter controls on the error log and audit viewer are keyboard-operable with programmatically associated labels and announced result counts on change.
- Announcement authoring is an ordinary accessible form: labelled fields, inline errors, error summary with focus management.

Success Outcome

Priya retrieves the complete hub-and-adaptor chain for a reported action from a single correlation ID in one filtered query, classifies the failure from one console in minutes, and contains it with an audited action reflected immediately in user navigation — while her own actions remain attributable to her and mission content remains outside her reach (JTBD-04.2, JTBD-04.3, JTBD-04.4; SM-20, PRD F11 acceptance signal).

Feature Touchpoints

Stage	Features
1. Start where the truth is	F11 , F16, F8
2. Confirm the symptom	F11, F16
3. Read the failures	F11 , F13, F9
4. Take the thread	F13, F10, F11
5. Read it as one story	F13 , F10
6. Classify and answer	F13, F11, F2

Stage	Features
7. Contain if needed	F11, F8, F13
8. Close the loop	F11, F13, F15

JRN-04.03: Rehearse Degradation on Purpose — Induce, Observe, Restore

Persona: PER-04 (Priya Raghunathan) **Scenario:** Before a demonstration — and, in a real program, before a production incident teaches it to her — Priya needs to verify that the platform degrades the way it claims to. She takes Investigation Management offline **on purpose** from the administrator-only failure-injection controls, watches the health monitor notice, watches the integration error log fill with correctly attributed entries, watches an investigator's queue degrade **visibly rather than blankly** in the next window, and then restores the spoke and watches everything clear automatically with nobody reloading and nobody re-authenticating. Today there is no way to induce a controlled failure at all, so resilience behavior is untested until it is production. **Related Jobs:** JTBD-04.2 (*health before complaint*), JTBD-04.4 (*resilience is rehearsable*); drives JTBD-01.4 and JTBD-02.4 from the other side of the glass **Entry Trigger:** Pre-demo rehearsal, or an evaluator asking "what happens when one of these is down?" **Demo Path:** □ **SECONDARY SCRIPTED DEMO — Segment 2b.** This is the **driving** half of the resilience segment: Priya acts here, the evaluator watches JRN-01.03 (investigator) and JRN-02.02 (adjudicator) in the other window. Always restore before moving to the next segment.

Preconditions (Seed Data Required)

#	Namespace	Required seed state
P1	hub	Failure-injection controls present, administrator-only, and demo-scoped (SCR-38), able to force any registered spoke into unavailable, slow, or erroring state and restore it.
P2	im	IM designated as the outage-demonstration spoke , with seeded items — including overdue ones — owned by the investigator persona, so removing it visibly changes counts on the observing screen.
P3	pvq	PVQ independently injectable, because the adjudicator-facing consequence (JRN-02.02) is a different and equally important demonstration.
P4	hub	A populated per-principal, per-source count cache so the degraded warning can quantify the omission rather than falling back to unquantified copy.
P5	hub	Health-probe interval and circuit-breaker thresholds tuned so degradation and recovery are each observable within a single check interval — a resilience demo with a ninety-second silence in the middle does not land.
P6	hub	An investigator session open in a second window before injection begins, so "no reload, no re-authentication" is observed rather than asserted.
P7	hub	Restore leaves no residue : no stuck circuit, no stale degraded flag, no orphaned error state. Verified as part of rehearsal, because a failure injection a restore does not fully clear is a firing criterion for this persona.

Journey Stages

Stage	Action	Touchpoint	Thinking	Feeling	Pain Point	Opportunity
1. Set the stage	Confirms all six applications healthy and the investigator's queue complete in the other window	SCR-24 system health + <i>(second window)</i> SCR-13	"Everything green, everything text-labelled. Baseline established."	Composed	Today: she has no baseline view at all; "healthy" is an informal belief per system	A single health surface that makes a baseline statable in one screenshot (F11)
2. Break it deliberately	Forces Investigation Management into unavailable from the failure-injection controls	SCR-38 failure-injection controls (administrator-only)	"Down it goes. Now let's see whether we told anyone."	Deliberate, slightly theatrical	Today: degradation cannot be rehearsed, so resilience behavior is discovered in production with an audience	Controlled, reversible failure injection — resilience becomes testable rather than hypothetical (F16)
3. Watch the monitor catch it	Watches health flip to unavailable within one check interval, with a transition time and text label	SCR-24 → IM row + check history	"Seven seconds. The monitor found it before any human could have."	Satisfied	Today: the first signal of a failing integration is an investigator calling to say his queue is empty	Detection by monitoring rather than by user report — the core of JTBD-04.2
4. Watch the log fill	Reads correctly attributed entries appearing in the integration-issue log with operation, error class, and correlation ID	SCR-25 integration issues	"Attributed to IM, named operation, real error class, and an ID I can follow."	Confident in the instrumentation	Today: failures leave traces in several places and nothing ties them together	A correctly attributed error-log entry within one health-check interval (PRD F11 acceptance signal)
5. Watch the user experience it	Switches to the investigator window: four sources still rendering, a named and quantified warning, IM-targeted actions disabled with an explanation, no error page anywhere	<i>(second window)</i> SCR-13 → degraded banner	"He can still work. He knows exactly what he's missing. Nobody had to call me."	This is the point	Today: the user sees an empty list, assumes the whole system is broken, and stops working	Degradation that is named, quantified, and survivable (SM-15, SM-16)
6. Put it back	Restores IM; watches health return to healthy, the warning clear, and the missing items reappear in the other window	SCR-38 restore → SCR-24 → <i>(second window)</i> SCR-13	"Cleared on its own. He didn't reload, and he never signed in again."	Assured, finished	Today: recovery is confirmed by asking users whether it looks better now	Automatic recovery via half-open probing, with no user action at all (SM-17)

System Response by Stage

Stage	System Response	Features
1. Set the stage	Health monitor reports current status, last successful check, latency, and recent history for every registered application, status conveyed by text and/or icon	F11, F16
2. Break it deliberately	Injection applied to the targeted spoke only; the control is administrator-only, server-authorized, and the injection action is itself audited	F16, F2, F13

Stage	System Response	Features
3. Watch the monitor catch it	Background probe fails; after the configured threshold the circuit opens and the hub stops hammering the spoke; health state transitions with a recorded time	F16 , F8
4. Watch the log fill	Adapter failures logged with timestamp, application, operation, error class, correlation ID, and affected principal, linked to the corresponding audit entries	F11, F13, F8
5. Watch the user experience it	Aggregate queue returns partial results with per-source status; the degraded banner names the application and quantifies the omission; actions targeting the unavailable spoke are pre-emptively disabled with a reason; the global error boundary guarantees no blank page and no stack trace anywhere	F16 , F5, F6, F14
6. Put it back	Half-open probing succeeds, the circuit closes, health returns to healthy, and open sessions receive restored data through polling — no reload, no re-authentication ; the warning clears through the same live region that raised it	F16 , F1, F5

Emotion curve (1 anxious → 5 confident): Baseline **4** → Inject **4** (*in control — but this is the moment she finds out whether the claims are real*) → Monitor catches it **5** → Log fills **5** → User's screen **5** (*named, quantified, still working — the payoff*) → Restore **5**.

Note on this curve. Priya's journey is the only one in this document that does not dip below 4. That is correct and deliberate: the anxiety of this scenario belongs to **Marcus and Dana**, in the other window, and the product's job is to keep their curves from collapsing. Priya's flat-confident curve *is* the outcome being demonstrated — an operator who is not surprised.

Key Moments

- **Demo Moment — Stage 5.** The window switch. An evaluator seeing "Investigation Management is unavailable — 12 items are not shown. The rest of your work is up to date." on a real user's screen, while the rest of that screen keeps working, is the entire resilience argument in one frame.
- **Decision Point — Stage 2.** Which spoke to break. IM is the designated outage spoke because its absence **visibly changes counts** for the investigator persona; injecting a spoke with no seeded work makes the demonstration abstract.
- **Delight Opportunity — Stage 6.** Automatic recovery with no user action. "He didn't reload, and he never signed in again" is the sentence to say out loud.
- **Risk — the whole journey.** Failure injection that leaves the platform in a state a restore does not fully clear. That is a firing criterion for this persona and a demo-day catastrophe: it would poison every segment that follows.
- **Risk — Stage 5.** A single error page, blank screen, or stack trace anywhere in the product during the injection fails the segment outright, regardless of how good the warning looks (SM-15, NFR-09).

Success Exit Criteria

- An induced adapter failure produces a **degraded/unavailable health state and a correctly attributed integration-error entry within one health-check interval** (PRD F11 acceptance signal).
- The forced outage produces a **visible, specific degraded warning across the product and no error page anywhere** (SM-15).
- Remaining sources render fully and **remain actionable** (SM-16).
- Restoring the spoke **clears the warning and restores data without user reload or re-authentication** (SM-17).
- The injection and restoration actions are **themselves audited** and attributable to Priya.

- The platform returns to a **clean baseline** — no stuck circuit, no stale flag — so the next demo segment runs on pristine behavior.

Failure and Alternate Paths

Condition	What Priya sees	Recovery
She injects slow rather than unavailable	Adapters time out at their configured bounds; the page stays interactive with loading states and then degrades to the named warning; one slow spoke never stalls the aggregate request	Restore; timeout policy verified
She injects erroring responses	Health reports degraded rather than unavailable; the error log records the error class distinctly; the circuit opens on repeated failure	Restore; error classification verified
Two spokes injected at once	Both named in the warning, each with its own count; the composition does not collapse into a generic message	Restore both
Recovery does not clear the warning	A defect, not an alternate path — automatic recovery is SM-17 and a blocking issue for the demo	Fix before demo
A user is mid-form when injection or recovery lands	The live region announces; focus is never moved ; entered data is untouched	Verified explicitly during rehearsal, because this is the most likely accessibility regression
Injection controls reachable by a non-administrator	A defect. The controls are administrator-only and server-authorized; a direct API attempt by any other role must be denied and audited	Covered by negative-path tests (F19)

Accessibility Notes

- **Keyboard-only and screen-reader completable end to end**, including the injection controls, the confirmation, and the restore.
- Failure-injection controls are **state-changing and consequential**: they need clear labels, an accessible confirmation with correct focus trapping and restoration, and an announced result — not a bare toggle whose effect is only visible somewhere else on the page.
- **Health and injected state are never colour-only**. "Unavailable (injected)" needs a text label; an operator must be able to tell an induced state from a real one, non-visually, so a demo state is never mistaken for a genuine outage.
- The observed degraded warning in the user window must be verified **with a screen reader during rehearsal**: announced via live region, **without stealing focus**, and also present in the static page structure for anyone arriving after the announcement.
- Check-history and error-log tables follow the standard accessible data-table patterns with announced result counts as new entries arrive.
- Transition times, correlation IDs, and error classes are selectable, copyable text.

Success Outcome

Priya induces a controlled spoke failure and observes the specific degraded warning appear across the product with **no error page anywhere**, a correctly attributed error-log entry within one health-check interval, and — on restore — the warning clearing automatically **without any user reloading or re-authenticating**: the demonstrable half of JTBD-04.2 and JTBD-04.4, and the enabling condition for JTBD-01.4 and JTBD-02.4 (SM-15, SM-16, SM-17).

Feature Touchpoints

Stage	Features
1. Set the stage	F11, F16, F8
2. Break it deliberately	F16 , F2, F13
3. Watch the monitor catch it	F16 , F11, F8
4. Watch the log fill	F11, F13, F9
5. Watch the user experience it	F16 , F5, F6, F14
6. Put it back	F16 , F1, F5, F11

Cross-Journey Patterns

Common Pain Points (the "today" costs that recur across personas)

These are the frictions that appeared in **three or more** journeys. Each is a single root cause with four different symptoms — which is precisely why a unified layer, rather than four better applications, is the right response.

#	Recurring pain (today)	Appears in	Single cause	Solved once by
X-01	Multiple logins for one piece of work , with mismatched session lifetimes so the system needed last has always expired	JRN-01.01, 01.02, 02.01, 03.01	Identity is owned per application	F0 + F1 — one session, honored by every adapter, for its whole life
X-02	Identifiers hand-carried between systems in a scratch document, with transposition as a recurring error class	JRN-01.01, 02.01, 03.01	No system accepts a reference from another	F6 related items + F7 — relationships resolved live by the hub
X-03	An empty screen that might be a lie — no way to distinguish "nothing here" from "the source could not be reached"	JRN-01.03, 02.02, 04.02, 04.03	Aggregation without per-source status	F16 — designed empty states and named, quantified degraded states as two different screens
X-04	Discovery is accidental — work, notices, and exceptions are found late or second-hand	JRN-01.02, 02.02, 03.01	Nothing computes exceptions across systems	F15 + F4 — server-side alert rules linking directly to the item that produced them
X-05	Confirmation that names nothing — "Saved" with no system, no state, no evidence — forcing manual re-verification	JRN-01.01, 02.01, 03.01	Writes are per-system and unattributed	F6 + F7 — confirmations naming system and resulting state, read back per spoke
X-06	Correlation by hand — reconstructing "what happened, in what order, across which systems" by merging timestamps	JRN-01.01, 02.01, 04.02	No correlation ID was ever issued	F13 — one correlation ID spanning an entire cross-system action

#	Recurring pain (today)	Appears in	Single cause	Solved once by
X-07	Role meaning differs per system , so nobody can say what they are actually entitled to and effort is invested in actions that turn out to be forbidden	JRN-01.02, 02.01, 03.02, 04.01	Authorization is implemented four times, four ways	F2 — one server-side policy engine, action sets computed per principal per resource
X-08	Context lost at every boundary — filters, sort, page, and the user's own train of thought	JRN-01.01, 01.02, 02.01, 03.01	Each application is a separate destination	F3 + F6 — one shell, breadcrumbs that carry cross-application context, restored queue state
X-09	Internal vocabulary leaking to people who cannot decode it	JRN-03.01, 03.02	Status display is an internal state dump	F4 + F14 — plain-language mapping applied server-side before anything reaches an applicant
X-10	Degradation cannot be rehearsed , so resilience behavior is first observed in production	JRN-01.03, 02.02, 04.03	No controlled failure mechanism exists	F16 failure injection — resilience becomes a test, not a hope

Shared Opportunities (built once, paid for by every persona)

- **One session, four experiences.** F0/F1 eliminate X-01 for every persona simultaneously. This is the single highest-leverage capability in the product, and it is invisible when it works — which is why the audit log's "exactly one authentication event" is the observable that proves it (SM-02).
- **Per-source status on every aggregate.** Every list, panel, and widget knows whether each source answered, failed, or answered with nothing. That one design decision produces the investigator's named warning (JRN-01.03), the adjudicator's refusal to decide blind (JRN-02.02), and the administrator's health picture (JRN-04.02) from the same substrate.
- **Server-computed action sets.** The same mechanism that keeps Dana from resolving an issue (JRN-02.01), Renée from touching another subject's record (JRN-03.02), and Priya from reading mission content (JRN-04.01/04.02) is one policy engine, evaluated per resource, on every request. **Three different demo moments, one implementation.**
- **The correlation ID as connective tissue.** Issued at the browser, propagated hub → adapter → spoke, landing in both the audit trail and the integration error log. It is what turns Marcus's flagship action into one chain (JRN-01.01), lets Priya answer a report in minutes (JRN-04.02), and gives every error state something a user can hand to an administrator.
- **Accessible data tables as a single pattern.** The work queue, the notices list, the audit viewer, the inventory, the health history, and the error log are the **same accessible table** with different columns. Building it once and well serves every persona; building it six times is how the console ends up non-conformant.
- **Designed empty states as a first-class deliverable.** Every journey has at least one path that ends in nothing to show. An empty state that explains and offers a next action is the difference between "the system is working and I have nothing to do" and "the system is broken."

Convergence Points (where journeys meet)

Convergence	Journeys	What happens	Why it matters to the demo
The same work item, two roles	JRN-01.01 ↔ JRN-02.01	Marcus resolves ISS-2207 ; Dana can read it and its narrative but is not offered the resolve action	The clearest live RBAC demonstration in the product — show them back to back
Producer → consumer of the record	JRN-01.01 → JRN-02.01	Marcus's disposition and narrative are the material Dana adjudicates weeks later, out of context	Makes the audit and narrative-quality requirements concrete rather than abstract

Convergence	Journeys	What happens	Why it matters to the demo
One outage, two consequences	JRN-04.03 → JRN-01.03 + JRN-02.02	Priya injects once; the investigator sees a named, quantified queue gap, the adjudicator sees a refusal to present an ambiguous empty list	One induced failure yields two persona-appropriate designed responses — the strongest evidence that resilience was designed, not patched
One registration, one user-visible consequence	JRN-04.01 → JRN-01.02	CVS registered in one window; its items appear in an open investigator session in the other	Extensibility is proven by the second window, not by the wizard
Mediated, never shared	JRN-01.01 ↔ JRN-03.01	Marcus acts on Renée's record; she never sees his findings, his notes, or that an issue was raised against her answer. Their only touchpoint is a request for information: a task to her, a pending item to him	A deliberate access-control boundary, not an oversight — and the setup for the zero-trust demonstration
Every action lands in the same trail	all ten	Authentications, traversals, writes, denials, adapter failures, console actions	The audit viewer is where every journey in this document can be seen to have happened

Cross-Cutting Accessibility Commitments

Stated once here because they apply to **every** journey and should be verified per journey rather than per screen:

1. **Every journey in this document is completable keyboard-only and with a screen reader, start to finish.** This is a completion criterion, not a note (SM-08).
2. **Cross-application traversal preserves focus and announces context.** The highest-risk step in the product (JRN-01.01, Stage 7) and the one a screen-reader user will notice first if it is wrong.
3. **Asynchronous updates — queue refreshes, action results, degraded warnings, new notices — are delivered via live regions and never steal focus,** with the single deliberate exception of a result the user is explicitly waiting for (the dual-system confirmation).
4. **Status is never colour-only,** anywhere, for anyone — including the administrator console, where a red/green health dot is both a conformance failure and an operational hazard (NFR-02).
5. **Source attribution lives in the accessible name,** not in a coloured badge. An adjudicator who cannot attribute a field cannot use it.
6. **Long, high-consequence forms** — issue resolution, determination, applicant submission, application registration — all require an error summary with focus management and in-page links to offending fields. A lost narrative is a lost user in four different ways.
7. **Usable at 200% zoom and 320px-equivalent width,** with the applicant surface (JRN-03.01) as the hardest case and the non-dismissible demo banner as a real layout constraint to test against (NFR-13, NFR-16).
8. **Denial, empty, degraded, and error pages are in scope for 508** — they are pages like any other, and they are where a user is already frustrated.

Journey-to-JTBD Traceability

Every journey stage below maps to the job it serves and the outcome that must be observable in a live synthetic-data demonstration. Stage names are abbreviated; full text is in each journey's stage table.

Journey Stage	JTBD ID	Expected Outcome
JRN-01.01:Sign in	JTBD-01.1	One authentication event serves the entire cross-system workflow; the audit log shows exactly one (SM-02)
JRN-01.01:Orient	JTBD-01.2	A newly raised PVQ issue on a case he owns is announced , not discovered late, and links directly to the item
JRN-01.01:Discover relationship	JTBD-01.1	The issue raised against a questionnaire answer is visible on the case , explained in words, resolved live rather than hard-coded
JRN-01.01:Traverse	JTBD-01.1	Traversal happens inside the same shell with case context carried in the breadcrumb — no interstitial, no new tab, no credential prompt
JRN-01.01:Resolve	JTBD-01.1	Disposition and narrative persist as a real action, re-authorized server-side at execution against both target resources
JRN-01.01:Dual confirmation	JTBD-01.1	Both spokes independently return the updated state through their own APIs; a partial write is never reported as success (SM-03)
JRN-01.01:Leave the record	JTBD-01.3	The whole action is retrievable as a single correlated chain , with exactly one audit record per state change (SM-19, SM-20)
JRN-01.02:Open the queue	JTBD-01.2	One queue spanning at least four of five spokes replaces three per-system lists; default view answers "what first" (SM-14)
JRN-01.02:Narrow	JTBD-01.2	Filtering and sorting are server-scoped, explainable across sources, and announce result count and sort state
JRN-01.02:Return from interruption	JTBD-01.2	Filters, sort, and page survive a 10–60 minute interruption without rebuilding
JRN-01.02:New source appears	JTBD-04.1	A newly registered application reaches an open session without sign-out or reload (SM-12)
JRN-01.03:Read the warning	JTBD-01.4	Degradation names the application and quantifies the gap ; never a blank screen and never an error page (SM-15, NFR-10)
JRN-01.03:Keep working	JTBD-01.4	Remaining sources render fully and remain actionable while one spoke is down (SM-16)
JRN-01.03:Blocked action	JTBD-01.4	Actions targeting an unavailable spoke are pre-emptively disabled with an explanation rather than failing mid-submission
JRN-01.03:Recovery	JTBD-01.4	The warning clears and data returns without reload or re-authentication (SM-17)
JRN-02.01:Read what was submitted	JTBD-02.1	The questionnaire as submitted is reachable in the same path as everything else bearing on the subject

Journey Stage	JTBD ID	Expected Outcome
JRN-02.01:Follow the issue	JTBD-02.1	Cross-system relationships resolve inline; traversal replaces four searches and four logins (SM-02, SM-04)
JRN-02.01:Read the narrative	JTBD-02.1	The investigator's disposition and narrative are visible, not merely that an issue existed and closed
JRN-02.01:Notice what she cannot do	JTBD-02.2	The same work item presents a different, server-computed action set per principal; a direct API call to the investigator-only action is denied and audited (SM-18)
JRN-02.01:Render determination	JTBD-02.2	Exactly one audit record, written before the success response, naming actor and the attributes held at the time (SM-19)
JRN-02.02:Ask the clock question	JTBD-02.3	Aging and blocked determinations are surfaced continuously, each alert resolving to a real populated destination (SM-05, SM-06)
JRN-02.02:Sort by age	JTBD-02.3	Date semantics are consistent across sources, so a sorted list is defensible
JRN-02.02:Empty issue list	JTBD-02.4	A genuine absence and an unreachable source are two different screens , distinguishable visually and non-visually
JRN-02.02:Read which one it is	JTBD-02.4	With PVQ offline, a specific named warning appears and no empty issue list is presented (SM-15, NFR-10)
JRN-03.01:Read where she stands	JTBD-03.1	"Where am I" answered within 30 seconds on a 320px viewport, as one process rather than two system statuses (SM-05, NFR-16)
JRN-03.01:Understand, not decode	JTBD-03.1	Zero internal system names, tier codes, or state abbreviations without plain-language explanation
JRN-03.01:See what she owes	JTBD-03.2	A short, ordered, unambiguous task list with due dates — not a systems inventory
JRN-03.01:Do the thing	JTBD-03.2	100% of tasks link directly to the action that discharges them ; nothing is lost to a session timeout (SM-06, NFR-14)
JRN-03.01:Confirmation	JTBD-03.2	The confirmation names what was received and what happens next, then appears in her own history
JRN-03.01:Read the notice	JTBD-03.3	Notices are delivered where status is checked, retained, with unread state, announced without stealing focus
JRN-03.02:Find the entry	JTBD-03.4	A self-scoped history produces evidence of a prior submission in under a minute
JRN-03.02:Check the edges	JTBD-03.4	Investigative content, issue items, designations, and case records are unreachable from every applicant surface

Journey Stage	JTBD ID	Expected Outcome
JRN-03.02:Wrong subject / wrong role	JTBD-03.4	Denied server-side , with a consistent non-enumerable error , and the denial written to the audit trail (SM-18)
JRN-03.02:Read the denials	JTBD-04.4	Denials — not only successes — are on the record, with actor, target, outcome, and correlation ID
JRN-04.01:Start the flow	JTBD-04.1	Onboarding is a guided configuration action performed in the UI, not an engineering project
JRN-04.01:Test before committing	JTBD-04.1	A live connection test gates submission; unreachable endpoints are caught before, not after
JRN-04.01:Let it describe itself	JTBD-04.1	Capabilities auto-discovered via <code>describe()</code> and confirmed, not transcribed from documentation
JRN-04.01:Review and submit	JTBD-04.4	Registration writes one audit record naming the administrator, the application, and the configuration
JRN-04.01:Watch it come alive	JTBD-04.1	Immediate appearance in inventory, health, navigation, and queue — under five minutes, zero code changes, zero restarts (SM-11)
JRN-04.01:Prove the consequence	JTBD-04.1	An investigator with an open session sees the new source's items without signing out (SM-12)
JRN-04.02:Start where the truth is	JTBD-04.2	Health, latency, and check history for every registered application are visible before a user reports a problem
JRN-04.02:Read the failures	JTBD-04.3	A filterable error log with application, operation, error class, correlation ID, and affected principal, linked to audit
JRN-04.02:Take the thread	JTBD-04.3	A correlation ID from a user-facing error state resolves the complete hub-and-adapter trace in one query (SM-20)
JRN-04.02:Read it as one story	JTBD-04.4	A cross-system action reads as one correlated chain — a record, not a reconstruction
JRN-04.02:Contain if needed	JTBD-04.3	Disabling an application is immediate in user navigation and queue, and is itself audited
JRN-04.02:Close the loop	JTBD-04.4	The administrator's own actions are attributable to her; mission content and audit mutation remain out of reach for everyone (NFR-07)
JRN-04.03:Break it deliberately	JTBD-04.4	Resilience is rehearsable: controlled, reversible, administrator-only failure injection, itself audited
JRN-04.03:Monitor catches it	JTBD-04.2	Degraded/unavailable state and a correctly attributed error-log entry appear within one health-check interval

Journey Stage	JTBD ID	Expected Outcome
JRN-04.03:Watch the user experience it	JTBD-01.4	A visible, specific degraded warning across the product and no error page anywhere (SM-15, SM-16)
JRN-04.03:Put it back	JTBD-04.2	Restoration returns the application to healthy with no manual intervention and no user reload or re-authentication (SM-17)

JTBD Coverage Check

All sixteen jobs are exercised by at least one journey; the five demo-critical jobs are exercised by a journey that is on the scripted path.

JTBD	Journeys	On demo path
JTBD-01.1 (<i>flagship</i>)	JRN-01.01	☐ Segment 1 (primary)
JTBD-01.2	JRN-01.02, JRN-01.01	☐ Segment 3b
JTBD-01.3	JRN-01.01, JRN-01.02	☐ Segment 1
JTBD-01.4	JRN-01.03, JRN-04.03	☐ Segment 2a
JTBD-02.1	JRN-02.01	☐ Segment 4
JTBD-02.2	JRN-02.01	☐ Segment 4
JTBD-02.3	JRN-02.02	Segment 2c
JTBD-02.4	JRN-02.02	Segment 2c
JTBD-03.1	JRN-03.01, JRN-03.02	☐ Segment 5a
JTBD-03.2	JRN-03.01	☐ Segment 5a
JTBD-03.3	JRN-03.01	Segment 5a
JTBD-03.4	JRN-03.02	☐ Segment 5b
JTBD-04.1	JRN-04.01, JRN-01.02	☐ Segment 3a
JTBD-04.2	JRN-04.02, JRN-04.03	☐ Segment 2b / 6
JTBD-04.3	JRN-04.02	☐ Segment 6
JTBD-04.4	JRN-04.02, JRN-04.03, JRN-03.02	☐ Segment 6

Journey-to-Feature Coverage Check

Feature	Journeys exercising it	Feature	Journeys exercising it
F0	01.01, 01.02, 01.03, 02.01, 03.01	F10	01.01, 02.01, 03.02, 04.02
F1	01.01, 01.02, 02.01, 03.01, 04.01, 04.03	F11	01.03, 04.01, 04.02, 04.03
F2	01.01, 01.02, 02.01, 03.01, 03.02, 04.01, 04.03	F12	01.02, 04.01
F3	01.01, 01.02, 03.01, 03.02	F13	01.01, 02.01, 02.02, 03.01, 03.02, 04.01, 04.02, 04.03
F4	01.01, 01.02, 02.02, 03.01, 03.02	F14	all ten journeys
F5	01.01, 01.02, 01.03, 02.01, 02.02, 03.01, 04.01, 04.03	F15	01.01, 01.02, 02.02, 03.01, 04.02
F6	01.01, 01.02, 01.03, 02.01, 02.02, 03.01, 03.02	F16	01.02, 01.03, 02.02, 04.01, 04.02, 04.03
F7	01.01, 02.01	F17	01.01, 01.02, 02.01, 03.01, 03.02, 04.01
F8	01.01, 01.02, 01.03, 02.02, 04.01, 04.02, 04.03	F18	the scripted demo path (all six segments)
F9	01.01, 01.03, 02.01, 02.02, 03.02, 04.01, 04.02	F19	03.02, 04.03 (negative-path and resilience assertions)

Gap note. F18 (demo operability) and F19 (test and accessibility verification) are exercised *by* these journeys rather than *within* them — they are the scaffolding that makes every journey repeatable and verifiable. Their acceptance evidence is the scripted demo path in the header and the per-journey Success Exit Criteria throughout.

Downstream Use

Consumer	What it takes from this document
STORY-MAP	Journey stages become the story-map backbone steps; per-journey Success Exit Criteria seed acceptance criteria
FRD / Functional Requirements	Stage → screen (SCR-xx) mapping and Failure/Alternate Path tables define the required states per screen, including the non-happy ones
Seed data (F17)	Every journey's Preconditions table is normative input to the seed corpus and to seed-time validation assertions
Demo script (F18)	The scripted demonstration path, its six segments, and the per-journey demo flags
Verification / Test Plan (F19)	Success Exit Criteria map to SM-01 ... SM-22; Accessibility Notes define the manual keyboard and screen-reader passes per journey
UX design	Emotion curves and Key Moments identify where to invest design effort — and where a dip is legitimate and should not be designed away